

Security-Aware Workload-State Counterfactual Verification and Nodal Net-Value Settlement for Spatially Coupled AI Data Centers

Anonymous Authors

Abstract—Artificial-intelligence (AI) data centers can provide spatial-temporal demand response (DR), but a meter-only baseline cannot determine whether a claimed reduction is supported by committed work or displaced by rebound. This uncertainty can turn a local meter decrease into an unverified network payment. This paper introduces an auditable verifier that binds scheduler-telemetry provenance to workload-feasible counterfactuals and security-aware settlement. A sparse cumulative linear program (LP) enforces release, deadline, conservation, capacity, migration, and cross-midnight completion. A second job-indexed LP checks each positive-energy job, while a validation-only risk envelope and pre-event information gate separate contract formation from locked execution truth. Signed nodal value and preventive N-1 constraints then bound payment over three held-out power-conversion scenarios. The 121-day trace-driven inventory uses 16 earlier days for selection and 54 locked days for evaluation. The complete-ledger verifier reaches normalized root-mean-square error (nRMSE) 0.342 with 1.726 MWh/day of mechanism-isolation false credit. A separately solved, gate-masked controlled event replay has zero baseline error, a 0.625-MW mean response-tracking error, and $F_1 = 1.000$ for the committed response policy; this is an information-boundary certificate, not a field treatment estimate. The indexed witness contains 12,293,445 admissible job-slot variables, has zero recorded job-energy residual, and yields 0.194 MWh of net event reduction under the declared runtime-plus-queue windows. A full 54-day RTS-24 panel evaluates all 37 finite non-islanding line outages in each of 864 frozen-profile cells, with a maximum post-contingency loading of 1.000000 p.u. The same indexed witness has zero job-to-profile residual before N-1 settlement and returns 0.004293 USD per representative interval. These results establish an observational, auditable benchmark for security-aware DR verification.

Index Terms—AI data centers, counterfactual verification, demand response, security-constrained dispatch, nodal settlement.

I. INTRODUCTION

The growth of large-model training and inference is converting data centers from nearly static industrial loads into large, software-defined grid resources. Interactive inference is latency constrained, elastic inference can be delayed briefly, and batch graphics-processing-unit (GPU) work can be deferred for hours or migrated between regions. These capabilities enable demand response (DR), but they also invalidate the usual interpretation of a meter decrease. A reduction at one facility can be followed by delayed consumption or simultaneous rebound at another node. Paying the positive local decrease can therefore reward a redistribution that creates little or negative network value.

Conventional DR measurement and verification estimates the unobserved no-event demand from historical meter data. Arithmetic historical baselines remain common in market practice under published market guidance [1], despite known bias and manipulation incentives [2]. Regression-based facility models quantify baseline uncertainty but do not encode a computing workload ledger [2]. More broadly, DR research establishes the economic role of responsive load, yet a meter-only estimator cannot determine whether queued jobs could actually have produced its predicted trajectory.

Data-center scheduling studies provide the missing physical state. Temporal deferral, dynamic capacity, and geographical load balancing can reduce energy cost and peak demand [3]. Geographical assignment provides the complementary

spatial control mechanism [4]. Recent work assesses data-driven flexibility from production traces and system-level security-constrained coordination, while balancing-market and clean-computing studies optimize spatial-temporal data-center dispatch [5], [6]. A recent power-system review also identifies quality-of-service (QoS), service-level-agreement (SLA), and market-participation constraints that are absent from meter-only accounting; clean-computing coordination provides a complementary dispatch context [7], [8]. Hierarchical optimal-consumption models further connect facility efficiency with response quality. Virtual-link market clearing and space-time remuneration establish how geographical and temporal transfers can enter electricity markets [9], [10]. Receding-horizon operation and uncertainty-aware frequency-regulation coordination further establish the need to retain future arrivals and multi-site state [11]. These methods optimize operator-controlled schedules; verification of an independently submitted DR counterfactual against an auditable workload ledger remains unresolved. This separation between computing feasibility and market verification motivates the auditable ledger used here. Recent production-trace studies quantify data-center flexibility and load-aggregator coordination, but leave this ledger-to-settlement coupling open [12], [13]. Recent whole-facility measurements likewise require a separate bottom-up conversion from GPU power to facility demand; the present study therefore keeps raw energy calibration, benchmark scaling, and network placement as separate evidence layers.

The second gap is electrical. Geographic workload control is commonly valued with energy prices or gross megawatts. Such accounting ignores the fact that nodal marginal value changes when congestion or a contingency constraint binds. Direct-current (DC) power flow and linear distribution factors provide transparent network sensitivities [14], [15]; N-1 dispatch can impose every credible post-contingency line limit through line-outage distribution factors (LODFs) [16]. No existing data-center DR verification framework jointly certifies workload feasibility, preserves signed spatial response, and settles the resulting security-aware network value. Alternating-current (AC) feasibility is a distinct nonlinear check, not a consequence of the DC certificate. Public workload and facility telemetry can support an observational replay of that ledger, but it does not constitute a causal utility event; a rigorous study must therefore separate measured-execution replay, counterfactual construction, and network-value verification to avoid conflating them with one field intervention.

This paper makes three power-system contributions. 1) This paper formulates a provenance-bound workload-state demand model in which release, deadline, conservation, cross-midnight completion, capacity, and migration constraints are enforced simultaneously. The counterfactual uses an exact aggregate energy-flow representation and a second globally solved job-indexed linear program (LP); the same immutable ledger is additionally replayed with measured contiguous job intervals, runtime, GPU count, and native power so that task-level feasibility is not inferred from a divisible-flow residual. The network-facing profile is then defined by aggregating the same indexed witness, and Experiment 22 independently

rechecks its job equalities, GPU nameplate bounds, and regional capacity rows before replaying that aggregation through the N-1 settlement model with a zero job-to-profile residual. 2) This paper constructs a validation-only, continuous-arrival counterfactual verifier with a pointwise credit envelope. The output is an attainable spatial-temporal load trajectory with constraints absent from an unconstrained meter-only prediction, and its statistical performance is evaluated under an explicit false-credit, under-credit, and contract-overpayment audit. 3) This paper develops signed space-time nodal remuneration and a bilateral net-value contract under preventive N-1 dispatch, and maps the feasible workload set into a payment uncertainty interval over held-out workload-to-power conversion. The central result is a coupling invariant that carries one committed workload state into physical verification, network valuation, and a set-valued settlement rule; the individual LP, dispatch, and allocation primitives are standard and are not claimed as new in isolation. Since public traces expose no utility geography, the spatial result is identified by a feature-stratified trace scenario and the complete permutation/penetration panel, not by an invented physical location. The certificate is conditional on trusted scheduler and telemetry provenance: a cryptographic digest verifies record integrity after commitment but does not authenticate semantically false fields.

The remainder of the paper is organized as follows. Section II defines the spatial workload and network settlement problem. Section III presents the verifier, security-aware value rule, and allocation mechanism. Section IV specifies the data, baselines, and locked protocol. Section V reports the core counterfactual, execution-replay, job-level, security, payment, continuous-horizon, and provenance experiments. Section VI concludes with the deployment boundary and future validation requirements.

II. SYSTEM MODEL AND PROBLEM FORMULATION

A. Spatial workload state

Let $\mathcal{S}$, $\mathcal{K}$, $\mathcal{D}$, and $\mathcal{T}$ denote source regions, workload classes, destination data centers, and 15-minute intervals; the symbol h is reserved for a calendar day in statistical panels. The measured ledger supplies arrival energy a_{skt} for source s , class k , and release interval t . The decision $x_{skd\tau} \geq 0$ is the energy served at destination d in interval τ . Its cumulative state is

$$y_{skt} = \sum_{\tau=0}^t \sum_{d \in \mathcal{D}} x_{skd\tau}. \quad (1)$$

With cumulative arrivals $A_{skt} = \sum_{\tau=0}^t a_{sk\tau}$ and deadline D_k , release and deadline feasibility are

$$y_{skt} \leq A_{skt}, \quad (2)$$

$$y_{skt} \geq A_{sk,t-D_k}, \quad t \geq D_k. \quad (3)$$

The state recursion equals service, terminal service equals total arrivals, and $\sum_{s,k} x_{skd\tau} \leq \bar{P}_d \Delta t$. Facility demand is

$$p_{dt} = P_d^{\text{fix}} + \Delta t^{-1} \sum_{s,k} x_{skd\tau}. \quad (4)$$

For rolling window end T^+ and final event-day release T^0 , the terminal state is not forced to consume unexpired future work. Instead,

$$y_{skT^+} = A_{sk, \max\{T^0, T^+ - D_k\}}, \quad (5)$$

which completes every job due inside the window and every event-day arrival, while retaining later real arrivals in the release constraints. The window starts $D_{\max} + 96$ intervals before the event day and ends $D_{\max}$ intervals after it; omitted earlier work must therefore expire before the evaluated day. Equations (2) and (3) prevent early execution and late completion, while (4) maps the feasible service state to the facility meter. The variable x is an aggregate, preemptive service envelope: it is valid for checkpointable batch work and does not by itself assert a nonpreemptive placement for every GPU job. For the measured MIT replay, each joined job is also retained

with its contiguous execution interval, runtime, GPU count, and native average power; Experiment 14 checks that witness independently. This two-level construction keeps counterfactual optimization sparse while making the task-level scope explicit. For the counterfactual task-level check, let $j \in \mathcal{J}$ index a joined job, r_j denote its submitted release, and let d_j be the service-window endpoint. Slurm “timelimit” is an allocation run-time declaration; it is converted to $R_j = \lceil \text{timelimit}_j / \Delta t \rceil$ slots and combined with a precommitted queue allowance $B = 96$ slots, so $d_j = r_j + R_j + B$. The unlimited Slurm sentinel is mapped to the precommitted horizon $H_\infty = 128$ slots before adding B ; finite declarations are retained exactly. The 0.001-MW-per-GPU nameplate is an audit-only feasibility precheck: if a declared window cannot serve E_j , the run fails closed and the window is not enlarged. Thus d_j is available at submission and never uses the observed completion time. Let E_j be measured energy, s_j its native region, and g_j its measured GPU count. The indexed service variable $u_{j\tau}$ is constrained by

$$\begin{aligned} u_{j\tau} &\geq 0, & \sum_{\tau=r_j}^{d_j-1} u_{j\tau} &= E_j, \\ u_{j\tau} &\leq g_j \bar{p}_{\text{GPU}} \Delta t, \\ \sum_{j: s_j=d, r_j \leq \tau < d_j} u_{j\tau} &\leq \bar{P}_d \Delta t, \\ \min_u \sum_{j,\tau} [w_{\text{batch}}(\tau - r_j) + r^{\text{DR}} \mathbf{1}\{\tau \in \mathcal{E}\}] u_{j\tau}. \end{aligned} \quad (6)$$

Here $\bar{p}_{\text{GPU}}$ is the precommitted 0.001 MW/GPU nameplate. Thus the counterfactual is a globally solved, job-indexed preemptive batch schedule with explicit release, submit-time service window, energy, site capacity, and GPU-count constraints. It is not obtained by replaying the measured profile; its event tariff is fixed before scoring. Experiment 19 solves (6) for all 71,128 positive-energy jobs. The observed $[t_j^{\text{start}}, t_j^{\text{end}}]$ interval is retained only to build the independent native replay used for scoring and never constrains $u_{j\tau}$. For the job-to-network certificate, the indexed witness is mapped without a second optimization. Let $\mathcal{J}_{skd}$ be the ledger partition whose source, class, and declared destination labels are s, k, d ; Experiment 19 uses the native assignment $d = s_j$ and therefore has no outcome-dependent migration:

$$x_{skd\tau}^{\text{job}} = \sum_{j \in \mathcal{J}_{skd}: r_j \leq \tau < d_j} u_{j\tau}, \quad (7)$$

$$p_{d\tau}^{\text{job}} = P_d^{\text{fix}} + \Delta t^{-1} X_{d\tau}^{\text{job}}, \quad X_{d\tau}^{\text{job}} = \sum_{s,k} x_{skd\tau}^{\text{job}}. \quad (8)$$

Aggregate migration remains a matched-workload control; (8) is the indexed replay profile. Experiment 22 reconstructs $X_{d\tau}^{\text{job}}$ from every stored $u_{j\tau}$, rechecks the indexed primal, and verifies equality before network valuation, so no separately optimized aggregate trajectory receives a network value. Workload deferral and geographic assignment follow [3], [4], while rolling-state control follows [11]; waiting and migration costs are $w_k(\tau - t)x_{skd\tau}$ and $m_{sd}x_{skd\tau}$, with service-time substitution differing only by a fixed-arrival constant.

For provenance, each retained job is represented by the canonical tuple and its digest in

$$\begin{aligned} \ell_j &= (\text{id}_j, t_j^{\text{submit}}, t_j^{\text{start}}, t_j^{\text{end}}, E_j, n_j^{\text{tel}}, q_j, \kappa_j, \sigma_j), \\ \mathcal{L} &= \text{sort}\{\ell_j : j \in \mathcal{J}\}, \\ h_{\mathcal{L}} &= \text{SHA256}(\text{serialize}(\mathcal{L})). \end{aligned} \quad (9)$$

where E_j is positive aggregated Data Center GPU Manager (DCGM) energy and n_j^{tel} is its telemetry-record count. The Secure Hash Standard [17], canonical ordering, one-to-one join, release-execution order, and raw-to-join energy equality are checked independently of the scheduling LP. Thus the flow witness certifies the committed ledger and measured

profile, while the digest only detects post-commit changes. Experiment 16 reconciles the measured envelope with the 118-MW nameplate fixed before the validation/test split.

The verifier treats scheduler/DCGM services as authoritative and assumes the digest is committed before the event decision. It cannot establish that a deadline, region, class, or energy field was truthful before signing; that requires upstream attestation and lies outside this optimization certificate.

B. Counterfactual, response, and false credit

For event set $\mathcal{E} \subset \mathcal{T}$, p^0 is the no-event counterfactual, $p^{1,\text{sim}}$ is the workload-optimized event trajectory, and p^{obs} is the independently measured MIT/DCGM execution trace. Because no utility event label is present, p^{obs} is an observational alignment target and never a causal outcome or pre-decision input. In mechanism-isolation rows $r = p^0 - p^{1,\text{sim}}$ is signed, so rebound is retained. For source $z \in \{\text{obs}, \text{sim}\}$, submitted and true credit are $\hat{r}^{(z)} = [\hat{p}^0 - p^{(z)}]_+$ and $r^{(z)} = [p^0 - p^{(z)}]_+$; source-specific false credit is

$$F^{(z)} = \Delta t \sum_{d,t \in \mathcal{E}} [\hat{r}_{dt}^{(z)} - r_{dt}^{(z)}]_+, \quad z \in \{\text{obs}, \text{sim}\}. \quad (10)$$

The source-specific quantity in (10) is evaluated separately for the simulated and observed trajectories. The simulated source defines the validation epigraph; the observed source is recomputed only in the locked replay. For nonnegative true credit, $p^{(\text{sim})} + r^{(\text{sim})} = \max\{p^0, p^{(\text{sim})}\}$, an audited identity rather than a new meter trajectory. The gross forecast credit is not a transfer: before the event, the operator freezes p^{con} , and after meter closure the payable response is the pointwise intersection

$$Q^{\text{pay}} = \Delta t \sum_{d,t \in \mathcal{E}} \min\{[\hat{p}_{dt}^0 - p_{dt}^{\text{obs}}]_+, [p_{dt}^{\text{con}} - p_{dt}^{\text{obs}}]_+\}. \quad (11)$$

Equation (11) uses only the submitted baseline, frozen contract, and closed meter. The trace-anchored p^0 is an offline over/underpayment diagnostic and never enters fitting, gate decisions, workload optimization, or transfer formation. We report gross, payable, oracle over/underpayment, nRMSE, bias, precision, recall, and F_1 .

C. Strategic reference scheduling

Consider an arithmetic mean of $H = 10$ reference days, event probability q , and DR price π^{DR} . If reference day j adds event-window service δR_j , the expected payment change is

$$\Delta \Pi = \frac{q\pi^{\text{DR}}}{H} \sum_{j=1}^H \delta R_j - \sum_{j=1}^H \Delta C_j. \quad (12)$$

Let c'_j be the right derivative of the minimum operating cost with respect to an event-window minimum-service constraint on day j . Linear-program sensitivity [18] gives the following exact local condition:

$$\frac{q\pi^{\text{DR}}}{H} > \min_j c'_j \iff \exists \text{ a strictly profitable local deviation.} \quad (13)$$

At a degenerate optimum the dual at the current service level can be zero even when the right-direction marginal cost is positive. Experiment 1 therefore certifies c'_j with two independent right finite differences of 5×10^{-4} and 10^{-3} MWh and a 10^{-6} USD/MWh agreement tolerance. The historical-baseline premise is documented in [2]; (12)–(13) are this paper's workload-specific specialization.

D. Network value

Let M be the data-center-to-bus incidence matrix and $p_t = p_t^{\text{native}} + Mp_t^{\text{dc}}$ the nodal demand vector. For generator-to-bus map C_g , generator bounds, and power-transfer distribution

factor (PTDF) H with line ratings $\bar{f}$, the piecewise-linear security-constrained economic dispatch (SCED) value is

$$V(p) = \min_g \sum_i C_i(g_i) \quad (14)$$

$$\text{s.t. } \mathbf{1}^\top g = \mathbf{1}^\top p, \quad g \leq g \leq \bar{g},$$

$$-\bar{f} \leq H(C_g g - p) \leq \bar{f}.$$

The incidence mapping aggregates device demand at its connection bus, and (14) is the standard lossless DC representation [14], [15]. The realized signed grid value is

$$\Delta V = V(p^0) - V(p^1). \quad (15)$$

Positive ΔV is a credit and negative ΔV a debit; no positive-only truncation is applied.

III. PROPOSED METHODOLOGY

A. Model-based framework and matched-information pre-estimation

The state, power, and network equations in Section II are the first stage of the methodology: they define the feasible workload state before any statistical target or payment is formed. Fig. 1 shows the complete chain. Meter history and the submitted-job ledger enter a statistical pre-estimator and the workload feasible set in parallel. The pre-estimator never receives execution energy used as scoring truth. For each event day, it uses only the information available under one of two declared protocols: lagged meter/calendar features for online baselines, or the complete submitted ledger for equally informed post-event comparisons. Ridge regression and gradient boosting are used as standard matched-information comparators [19], [20]; extremely randomized trees and synthetic control provide two additional nonlinear and donor-pool comparators [21], [22]. Two stronger comparators are also solved: median-loss gradient boosting with the complete submitted ledger and a nonnegative synthetic control whose weights globally minimize pre-event absolute error over 28 donor days. To separate the effect of projection from that of pre-estimation, the complete-ledger quantile output is also passed through the same exact workload-feasible projection, with its penalty chosen by contiguous validation folds. The same complete ledger is supplied to the post-event learner, the selected single projection, and the proposed convex verifier.

B. Sparse cumulative workload-state verifier

For statistical estimate $\tilde{p}$ and each predeclared penalty ρ_ℓ , the first-stage program solves

$$\min_{x,y,z} C_w(x) + \rho_\ell \sum_{d,t} z_{dt} \quad \text{s.t.} \quad z_{dt} \geq \pm(p_{dt}(x) - \tilde{p}_{dt}), \quad (16)$$

subject to (1)–(4). Absolute-value epigraphs are standard convex constructions [18]; using them to project a meter baseline onto an auditable workload ledger is specific to this work. Here C_w denotes workload operating cost; C_i denotes generator-segment cost in (14), and $V(p)$ denotes the network dispatch value. Thus (16) changes the objective, not the physical feasible set. Six globally solved schedules $p^{(\ell)}$ share identical arrivals and constraints. Their validation-only convex combination $\tilde{p} = \sum_\ell \alpha_\ell p^{(\ell)}$, $\alpha \geq 0$, $\mathbf{1}^\top \alpha = 1$, remains feasible because the common scheduling set is convex.

The coefficients minimize validation squared error subject to both total and daily-tail mechanism-isolation false-credit budgets, where conditional value at risk (CVaR) measures the latter [23]. Let ℓ^* be the single candidate with minimum worst contiguous-fold nRMSE and define $B_j^{\text{cap}} = F_j^{(\text{sim})}(\ell^*)$, where $F_j^{(\text{sim})}$ is obtained from the explicit submitted-minus-true




Fig. 1. Verification-to-settlement framework. Validation labels select the counterfactual risk contract; locked execution truth is used only for scoring. The right branch computes aggregate N-1 value and exact participant shares.

credit definition above. The observed-meter source is audited separately after the locked decision. For reserve $\beta \in (0, 1]$,

$$\sum_j F_j^{(\text{sim})}(\alpha) \leq \beta \sum_j B_j^{\text{cap}}, \quad (17)$$

$$\text{CVaR}_{0.75}(F_j^{(\text{sim})}(\alpha)) \leq \beta \text{CVaR}_{0.75}(B_j^{\text{cap}}). \quad (18)$$

For implementation, the positive part is not differentiated through a sort. For each event sample n and validation day j , the program introduces $f_n \geq 0$, a CVaR threshold $\nu \geq 0$, and tail slacks $u_j \geq 0$:

$$f_n \geq p_n^\alpha - p_n^{(\text{sim})} - r_n^{(\text{sim})}, \quad (19)$$

$$F_j^{(\text{sim})} = \sum_{n \in j} f_n, \quad (20)$$

$$F_j - \nu - u_j \leq 0, \quad (21)$$

$$\nu + \frac{1}{0.25J} \sum_j u_j \leq \beta \text{CVaR}_{0.75}(B_j^{\text{cap}}). \quad (22)$$

Here $r_n^{(\text{sim})} = [p_n^0 - p_n^{(\text{sim})}]_+$ is the true mechanism-isolation credit. The equivalent threshold identity is checked in the implementation, but the two credit terms remain explicit in the model. This epigraph is never substituted for the frozen contract profile p^{con} in (11). The explicit epigraph in (22) is the implementation used in all reported fits. Together with the simplex equality, box bounds on α , and the total budget in (17), these are linear constraints and the squared error objective is a convex quadratic program. HiGHS first supplies a feasible epigraph point; the full epigraph is then solved with a sparse primal-dual trust-region method. The returned stationarity residual, every linear epigraph residual, the simplex equality, and all bound residuals are recomputed from the final primal point. The optimizer status and these residuals are stored for every fit. Four contiguous folds select β using validation nRMSE, then risk ratio; held-out noninferiority is reported as an outcome and is not used as a hidden constraint. A second exact projection targets $\bar{p}$ while imposing, for every event sample,

$$p_{dt}^{\text{safe}} \leq p_{dt}^{(\ell^*)}. \quad (23)$$

The pointwise cap in (23) is the upper side of the contractual credit band; it is evaluated on every event slot before settlement.

a) *Two-sided credit band.*: The upper envelope alone would admit an arbitrarily conservative zero-credit profile. We

therefore impose a lower band tied to the independently fitted convex target. Define the componentwise risk floor $p_{dt}^{\text{floor}} = \min\{p_{dt}^{(\ell^*)}, \bar{p}_{dt}\}$, where $\bar{p}$ is the validation-fitted convex target before the final workload projection. The event-slot constraints are

$$p_{dt}^{\text{floor}} - \varepsilon \leq p_{dt}^{\text{safe}} \leq p_{dt}^{(\ell^*)}, \quad (d, t) \in \mathcal{D} \times \mathcal{E}, \quad (24)$$

where $\varepsilon = 1$ MW is fixed before the locked test split. The lower bound is implemented together with the workload constraints, not as a post-processing clip. Thus the output remains an attainable schedule and can move below the single reference when the independent risk fit supports a safer trajectory. The total constraint (17) controls aggregate exposure and (18) controls the worst daily tail. The upper side is the deterministic payment cap; the lower side is a declared under-credit regularizer relative to p^{floor} , not a distribution-free guarantee against an arbitrary unseen meter trace. Overpayment against the realized event is handled by the contract-capped settlement rule (11). The selected single feasible projection proves nonemptiness, and the locked results report both margins and the post-event payable amount.

b) *Decision-time information boundary.*: The ex-post selector may use the complete ledger, but a gate deployment removes all post-gate arrivals before target construction and optimization. The target is refit with the matched-information metadata model on this masked ledger. A no-tariff LP freezes the contract baseline and commits 5 MWh of event service; a second LP applies the positive DR price to the same ledger and carries unexpired batch state to the next checkpoint. Reserved future jobs cannot be credited. Baseline error is computed only for the submitted contract; the tariff-bearing plan is scored as planned reduction against the closed meter and then settled by (11), keeping the two roles distinct.

To separate capacity planning from payment eligibility, the gate publishes a causal reserve envelope. For candidate quantile η and event day d ,

$$\hat{a}_{sk\tau}^{\text{res}}(d; \eta) = Q_\eta(\{a_{sk\tau}(h) : h < d\}), \quad \tau \geq t_{\text{gate}}, \quad (25)$$

Equation (25) uses temporally earlier days only. Candidates are selected on validation days under a declared false-credit budget and used only for capacity planning; no reserved job enters payment before ledger commitment.

The cumulative state reduces the constraint-matrix nonzeros from $O(|\mathcal{S}||\mathcal{K}||\mathcal{D}||\mathcal{T}|^2)$ for explicit prefixes to $O(|\mathcal{S}||\mathcal{K}||\mathcal{D}||\mathcal{T}|)$. A 512-slot deadline is handled by (5) with every observed future arrival retained; no zero-arrival comple-

tion buffer is used in the final continuous-horizon certificate. The day trajectory is embedded by a lexicographic pair of linear programs: minimum L1 distance first, followed by minimum operating cost on that optimal-distance face. Hence no penalty weight trades physical fit against cost. All linear programs are solved to global optimality with HiGHS [24]. The resulting state-to-credit sequence is summarized in Fig. 2; the execution replay and interval audit are reported separately in Experiments 13–15.

C. Signed nodal and N–1 settlement

Let $\lambda^0 \in \partial V(p^0)$ be the nodal demand-value subgradient. Signed linear settlement is $(\lambda^0)^\top(p^0 - p^1)$, while (15) is the realized system-value benchmark and the payoff of an explicit bilateral avoided-cost contract; it is not relabeled as an independent system operator (ISO) market payment. The convex subgradient inequality [18] yields the global certificate

$$0 \leq (\lambda^0)^\top(p^0 - p^1) - \Delta V = V(p^1) - V(p^0) - (\lambda^0)^\top(p^1 - p^0). \quad (26)$$

Thus the linear rule globally upper-bounds exact value in the implemented polyhedral market model. Equality holds only when both endpoints share a supporting dual subgradient.

For market-compatible remuneration, let $\mathcal{W}$ contain the complete response cycle from the earliest deadline-feasible anticipatory shift through the latest recovery interval. The space–time rule is

$$\Pi^{\text{ST}} = \Delta t \sum_{t \in \mathcal{W}} \sum_d \lambda_{dt}^0 (p_{dt}^0 - p_{dt}^1). \quad (27)$$

This is the nodal virtual-link remuneration structure of [9], [10], applied here to a verified counterfactual. Summing site payments reproduces (27) exactly.

Space–time energy remuneration is only one component of the contracted DR product. Let $\mathcal{D}^{\text{part}}$ be the predeclared participating sites, $\mathcal{E}$ the event intervals, and r^{DR} the cleared service price. The signed delivered capacity service and total operator value are

$$Q^{\text{cap}} = \Delta t \sum_{d \in \mathcal{D}^{\text{part}}} \sum_{t \in \mathcal{E}} (p_{dt}^0 - p_{dt}^1), \quad (28)$$

$$G = r^{\text{DR}} Q^{\text{cap}} + \Pi^{\text{ST}}, \quad (29)$$

where p^0 is the minimum-cost no-event schedule under the identical continuous arrival window. Capacity credits and ex-post participation constraints follow established DR mechanism design [25], [26]; the signed Π^{ST} term separately debits spatial transfer and recovery. Equations (28)–(29) jointly define the signed service and operator value.

Let x^1 and x^0 denote the workload service schedules inducing p^1 and p^0 , respectively. Define $C_{\text{opp}} = C_w(x^1) - C_w(x^0)$ as the participant opportunity cost and $S = G - C_{\text{opp}}$ as the transferable transaction surplus. With disagreement utilities equal to zero, the symmetric Nash solution [27] activates the contract iff $S \geq 0$ and sets

$$z = \mathbb{I}\{S \geq 0\}, \quad P^{\text{B}} = z \left(C_{\text{opp}} + \frac{S}{2} \right), \quad u_{\text{dc}} = u_{\text{op}} = z \frac{S}{2}. \quad (30)$$

Thus (30) makes voluntary nonactivation the explicit outside option. For every activated contract, both utilities are nonnegative and the single transfer P^{B} is exactly budget balanced. The bilateral transfer is reported separately from both ISO nodal remuneration and the exact dispatch value benchmark. Unlike event-only accounting, $\mathcal{W}$ debits both pre-event load increases and delayed rebound; workload conservation gives $\sum_{d,t \in \mathcal{W}} (p_{dt}^0 - p_{dt}^1) \Delta t = 0$ up to solver tolerance.

Security-aware settlement augments (14) with every finite non-islanding single-line outage. If $L_{\ell k}$ is the LODF from outage k to monitored line ℓ , then

$$-\bar{f}_\ell \leq (H_\ell + L_{\ell k} H_k)(C_{g\ell} - p) \leq \bar{f}_\ell, \quad \ell \neq k. \quad (31)$$

Equation (31) is the standard preventive N–1 LODF formulation [16]; the experiment imposes all credible rows simultaneously, with no contingency screened out. Native public ratings are retained. Islanding outages are reported and excluded because a connected single-reference PTDf cannot represent separate islands, consistent with the scope of DC security models and the N–1 planning criterion [28].

D. Scenario-robust N–1 payment certificate

The pointwise envelope controls credited energy but does not by itself certify network payment because spatially different demand vectors can activate different security constraints. Let $\mathcal{P} = \text{conv}\{p^{(1)}, \dots, p^{(L)}\}$ be the convex hull of complete workload-feasible schedules and let $p^{\text{cap}} \in \mathcal{P}$ be the validation-selected single feasible projection. Its penalty ρ_{ℓ^*} is selected by contiguous validation and is reported separately from the payment-target penalty. The feasible-quantile projection is retained as an external matched transfer comparator in Section V. Let Ξ_3 contain the 10th, 50th, and 90th percentiles of held-out per-job measured-to-predicted GPU energy. These predeclared empirical scenarios define the finite power-conversion contract supported by independent telemetry [29]. For each settlement day, the payment-certified verifier defines $B_\xi^{\text{cap}} = \Delta t \sum_{t \in \mathcal{E}} V_{N-1}(p_t^{\text{native}} + M\xi p_t^{\text{cap}})$ and solves

$$\begin{aligned} \text{lexmin}_{\alpha, e, \eta, \{g_t^\xi\}} \quad & \left(\sum_{d,t \in \mathcal{E}} e_{dt}, -\eta \right) \\ \text{s.t.} \quad & p^\alpha = \sum_{\ell=1}^L \alpha_\ell p^{(\ell)}, \\ & \alpha \geq 0, \quad \mathbf{1}^\top \alpha = 1, \quad 0 \leq \eta \leq 1, \\ & e_{dt} \geq \pm(p_{dt}^\alpha - p_{dt}^{\text{safe}}), \\ & g_t^\xi \text{ satisfies (14) and (31),} \\ & \quad \text{with } p_t = p_t^{\text{native}} + M\xi p_t^\alpha, \\ & \quad \quad t \in \mathcal{E}, \quad \xi \in \Xi_3, \\ & \Delta t \sum_{t \in \mathcal{E}} \sum_i C_i(g_{it}^\xi) \\ & \quad \leq (1 - \eta) B_\xi^{\text{cap}}, \quad \xi \in \Xi_3. \end{aligned} \quad (32)$$

Generator costs use the same declared piecewise-linear epigraph as market settlement. Two global linear programs implement the lexicographic objective: the first obtains the minimum trajectory error, and the second preserves that optimum within solver tolerance while maximizing the common fractional margin η . Existence of the primal dispatches in (32) is equivalent to the corresponding upper bound on the optimal N–1 value; the unit weight on p^{cap} proves nonemptiness in every scenario. Consequently, for any $\xi \in \Xi_3$ and realized event load $p^{1,\xi}$, daily payment obeys

$$\begin{aligned} & \sum_{t \in \mathcal{E}} \Delta t [V_{N-1}(p_t^{\text{native}} + M\xi p_t^\alpha) - V_{N-1}(p_t^{1,\xi})] \\ & \leq \sum_{t \in \mathcal{E}} \Delta t [V_{N-1}(p_t^{\text{native}} + M\xi p_t^{\text{cap}}) - V_{N-1}(p_t^{1,\xi})]. \end{aligned} \quad (33)$$

The common realized-cost term cancels, so the certificate requires no execution truth and holds samplewise. Both programs contain every finite non-islanding contingency. The formulation follows the standard SCED primal and linear-program epigraph foundations [15], [18]. Here $\mathcal{P}$ contains the six independently solved first-stage projections from (16) and the matched feasible-quantile projection. The risk-constrained verifier is an external target in the absolute-deviation objective and is not a selectable candidate; the quantile schedule is used only as an independent transfer comparator, not as the contractual payment cap.




Fig. 2. Sparse cumulative workload-state verifier. Real release-to-deadline edges feed the global feasibility LP; the pointwise credit envelope and the finite/interval payment gates are applied only after the workload state is certified.

The finite cap is complemented by a set-valued payment certificate. The contractual cap profile is denoted $p^{\text{cap}} = p^{(\ell^*)}$, where ℓ^* is selected by contiguous validation nRMSE; the independent payment-target profile p^{pay} is selected separately by validation payment error and is used only as the target of the certificate. For a realized counterfactual value, define $\Delta V_\xi(p) = V_{N-1}(p, \xi) - V_{N-1}(p^{1, \xi})$ and the continuous segment $p_\theta = (1 - \theta)p^{\text{cap}} + \theta p^{\text{cert}}$, $0 \leq \theta \leq 1$. The induced interval is

$$\mathcal{I}_\xi = \left[\min_{0 \leq \theta \leq 1} \Delta V_\xi(p_\theta), \max\{\Delta V_\xi(p^{\text{cap}}), \Delta V_\xi(p^{\text{cert}})\} \right]. \quad (34)$$

The lower endpoint is obtained by one exact joint N-1 SCED LP with a shared segment parameter; the upper endpoint follows from convexity of the optimal dispatch value. The selected cap payment is an element of $\mathcal{I}_\xi$. The endpoints use only validation-frozen feasible profiles and the realized counterfactual needed for settlement; no oracle baseline selects or certifies the interval. Interval width is reported as uncertainty cost, while the seven-profile hull remains the finite-scenario payment certificate in (32).

E. Participant allocation

For participants N , define $v(S)$ as signed dispatch value when coalition S moves from baseline to actual demand and nonmembers remain at baseline. The exact allocation is

$$\phi_i(v) = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(|N| - |S| - 1)!}{|N|!} [v(S \cup \{i\}) - v(S)]. \quad (35)$$

This is the Shapley value [30]; the characteristic function based on signed grid value is the proposed specialization. Grouping (35) by coalition size cancels every intermediate value and gives $\sum_i \phi_i = v(N)$ when $v(\emptyset) = 0$. We enumerate all 16 coalitions for four sites and all 256 coalitions for eight contractual portfolios; no permutation estimator is used. For larger auditable contracts, exchangeable slices at the same electrical site are grouped by their member count. The exact sum retains the binomial multiplicity of every labeled coalition and evaluates only $\prod_g (m_g + 1)$ distinct count states. This symmetry reduction remains the Shapley value and is evaluated through 20 participants.

F. Analytical guarantees

Proposition 1 (manipulation boundary): Under a ten-day arithmetic baseline, differentiable one-sided optimal costs, and an interior feasible increase in at least one reference-day event service, a strictly profitable local manipulation exists if and only

if (13) holds. To prove the statement, differentiate (12) along any feasible reference-day direction. The marginal payment is the same $q\pi^{\text{DR}}/H$ for every day, whereas the least physical cost is $\min_j c'_j$. A strict excess selects a positive direction on an argmin day; if no strict excess exists, every nonnegative local direction has nonpositive first-order profit. Equality admits indifference but does not imply strict profit. Experiment 1 computes c'_j from the honest program's right-hand-side dual and solves the strategic program independently, preventing the behavioral result from being reused as its own certificate.

Proposition 2 (workload and two-sided credit feasibility): For fixed arrivals, every convex combination of solutions to (1)–(4) is feasible, and the envelope output satisfies the deterministic two-sided band (24) on every event cell. Consequently, for every possible unseen execution c_{dt} , its false-credit energy is no larger than that of the selected cap $p^{(\ell^*)}$, while its additional under-credit is bounded by $\varepsilon|D||\mathcal{E}|\Delta t$ relative to the independent floor p^{floor} . Affine equalities are preserved by the simplex coefficients; each linear inequality is preserved by convexity. The upper inequality in (24) and monotonicity of $[z - c]_+$ give

$$[p_{dt}^{\text{safe}} - c_{dt}]_+ \leq [p_{dt}^{(\ell^*)} - c_{dt}]_+.$$

The lower inequality gives $[c_{dt} - p_{dt}^{\text{safe}}]_+ \leq [c_{dt} - p_{dt}^{\text{floor}}]_+ + \varepsilon$. Summation over sites and event intervals proves both bounds. The result is a feasibility-and-band certificate determined by the declared constraints and ε , not by the distribution of the locked test set.

Proposition 3 (security-aware value consistency): For the piecewise-linear base-case or N-1 dispatch, exact signed settlement equals the corresponding optimal-value difference and the nodal linear rule upper-bounds it by (26). The feasible-set right-hand side is affine in nodal demand, so the optimal value is convex and piecewise affine. Applying the subgradient inequality at p^0 gives (26); applying it again at p^1 yields

$$0 \leq \Delta V_{\text{lin}} - \Delta V \leq (\lambda^1 - \lambda^0)^\top (p^1 - p^0) \leq \|\lambda^1 - \lambda^0\|_* \|p^1 - p^0\|.$$

Here $\|\cdot\|$ and $\|\cdot\|_*$ denote a fixed dual pair of norms on the nodal-demand and price spaces. The same proof applies after stacking all rows in (31), because they remain linear. Gross positive-only settlement has no analogous identity: truncation changes the response vector before it enters the value function.

Proposition 4 (robust daily payment noninferiority): Every solution of (32) is workload feasible and its exact N-1 payment is no larger than the contractual cap payment for every realized event trajectory and every declared $\xi \in \Xi_3$. Workload feasibility follows from Proposition 2. For each ξ , the embedded dispatch is a feasible witness whose cost satisfies the corresponding last constraint in (32); optimal N-1 dispatch can only have lower

cost. Subtracting the same scenario-specific realized cost gives (33). Conditional on the telemetry-derived finite set Ξ_3 , the proof is deterministic and contains no test-outcome assumption.

Proposition 5 (continuous worst-case interval certificate): For a fixed workload profile p , the optimal piecewise-linear N-1 dispatch value $V_{N-1}(p; \xi)$ is convex in a scalar conversion factor $\xi \in [\xi_L, \xi_U]$. Therefore

$$\max_{\xi \in [\xi_L, \xi_U]} V_{N-1}(p; \xi) = \max\{V_{N-1}(p; \xi_L), V_{N-1}(p; \xi_U)\}.$$

The value function is the supremum of affine dual objectives in ξ , hence convex; a convex function on a compact interval attains its maximum at an endpoint. If the two endpoint inequalities $V_{N-1}(p^\alpha; \xi_e) \leq V_{N-1}(p^{\text{cap}}; \xi_e)$ hold for $e \in \{L, U\}$, then the certified candidate also has no larger worst-case interval cost than the contractual cap. This interval result is a worst-case bound; pointwise dominance at every interior ξ is asserted only for the finite set Ξ_3 in Proposition 4. Experiment 15 evaluates the capacity-safe endpoints with the exact four-segment contractual N-1 model, while Experiment 9 independently replays Ξ_3 at 40 segments; the two resolutions are not conflated.

Proposition 6 (workload-segment payment interval): For a fixed conversion factor and realized counterfactual, the lower endpoint in (34) is the exact minimum payment over the declared continuous segment because the shared segment parameter and all interval-wise preventive N-1 dispatch variables are solved in one joint LP. The upper endpoint is the exact maximum because the optimal dispatch value is convex in the affine profile and a convex function attains its maximum on a compact segment at a vertex. Thus (34) is exact for the contractual segment, not a grid approximation or a claim that the segment contains every physically possible schedule. Experiment 15 reports the joint-LP minimum, endpoint maximum, and resulting width at both held-out conversion endpoints; no oracle baseline selects the segment.

Proposition 7 (provenance-bound workload witness): Suppose the canonical tuple in (9) is computed from the retained scheduler-DCGM join after the declared duplicate and positive-energy rules. If the certificate verifies one-to-one joining, release-execution ordering, raw-to-join energy conservation, and zero LP residuals, then the job-level flow witness is a feasible representation of that committed ledger. Under the collision-resistance assumption of SHA-256, any post-commit modification of a tuple or its canonical ordering changes the digest except with negligible probability. The statement binds the computational witness to the committed records; it does not certify that an uncommitted pre-event submission was truthful. Experiment 16 checks each premise independently and records the raw source hashes, so the provenance claim is separate from the scheduling solver's own feasibility status.

Proposition 8 (indexed counterfactual feasibility): If the job-indexed program (6) returns an optimal solution, every retained job receives exactly its committed energy within its submitted release-deadline window, every site-slot total respects the declared capacity, and each service interval respects the measured GPU-count power bound. These statements follow directly from the equality, bound, and capacity rows of (6); they do not rely on aggregate-flow rounding. The guarantee is for preemptive checkpointable service. Nonpreemptive contiguous execution is asserted only for the independently observed intervals audited in Experiment 14.

Proposition 9 (typed ledger-to-settlement coupling invariant): Let $u \in \mathbb{R}^{|\mathcal{J}| \times |\mathcal{T}|}$ be the indexed service array, let $B_{dj} = \mathbf{1}\{s_j = d\}$ be the committed job-to-region incidence, and let $M \in \{0, 1\}^{|\mathcal{B}| \times |\mathcal{D}|}$ map regions to network buses. Define the

TABLE I
COMPLETE DATA AND EVALUATION PROTOCOL.

Item	Declared use and count
BurstGPT	5,188,507 requests; 2.176 billion tokens; inference arrivals and observed service
MIT scheduler	287,173 raw rows; 104 duplicate rows removed, leaving 287,069 canonical rows; submission time, declared timelimit, GPU count, and queue state
MIT GPU telemetry	96,893 rows; 68,664 tensor-window jobs and 71,128 full-horizon positive-energy jobs; measured execution energy
Time resolution	15 min; 121 calendar days; 121 valid complete days
Evaluation	16 validation days followed by 54 untouched test days
Networks	IEEE Reliability Test System 24-bus (RTS-24), IEEE-30, IEEE-39, and IEEE-118 for AC/N-1 panels; IEEE-300 is added in the resolution-convergence panel
Provenance	Immutable scheduler/DCGM join digest, raw-to-join energy conservation, and measured regional capacity envelope

certificate residuals

$$\begin{aligned} r_j^{\text{job}} &= \sum_{\tau=r_j}^{d_j-1} u_{j\tau} - E_j, \\ r_{d\tau}^{\text{agg}} &= X_{d\tau} - \sum_j B_{dj} u_{j\tau}, \\ r_{b\tau}^{\text{map}} &= p_{b\tau} - P_{b\tau}^{\text{fix}} - \Delta t^{-1} \sum_d M_{bd} X_{d\tau}. \end{aligned}$$

If the typed dimensions are satisfied and $\|r_j^{\text{job}}\|_\infty, \|r_{d\tau}^{\text{agg}}\|_\infty, \|r_{b\tau}^{\text{map}}\|_\infty \leq \varepsilon$, then the profile used by the security-constrained dispatch is the image of the same indexed witness up to ε in the declared units. At $\varepsilon = 0$, an independently optimized aggregate trajectory cannot receive a network value. Signed value, payment caps, and Shapley allocations therefore act on one committed p , with $\sum_i \phi_i = v(N)$. Experiment 22 evaluates all three residual arrays, the GPU and site-capacity bounds, and the network dispatch only after the checks; the ledger, conversion, and baseline remain committed between verification and settlement.

IV. EXPERIMENTAL DESIGN

A. Data, preprocessing, and information boundary

Table I reports the complete data inventory. BurstGPT contributes 5,188,507 successful requests and observed token service [31]. MIT Supercloud contributes scheduler submissions and independent Data Center GPU Manager energy. The complete positive-energy scheduler-DCGM join contains 71,128 jobs, of which 68,664 start within the common 121-day tensor window used by Experiments 1–13 [29]. Submission time creates queue arrivals; measured execution intervals create scoring truth. A held-out GPU-power regression uses 28,413 observations and obtains $R^2 = 0.899$, mean absolute error (MAE) 11.22 W, and root-mean-square error (RMSE) 15.82 W. All 99th-percentile workload scaling factors are fitted on the first 40 complete days of the training portion; validation and locked-test days are excluded from that fit. Per-job measured-to-predicted energy ratios define Ξ_3 in (32); no synthetic row replaces an invalid record. The 68,664 count is the 121-day workload tensor used by Experiments 1–13. Experiment 14 does not truncate at that tensor boundary; it extends the observed replay horizon to the last scheduler completion and retains all 71,128 positive-energy joined jobs for its independent job-level flow witness. Experiment 19 uses the same join but replaces the observed completion interval with the submit-time allocation-runtime window plus a precommitted 96-slot queue allowance; the unlimited sentinel receives the same rule before the allowance is added. This is a contract-window construction, not a submit-to-completion interpretation of Slurm “timelimit”.

The public traces do not contain utility bus locations, so the four-region coupling is evaluated under a trace-driven

benchmark setting. Regions are assigned by a predeclared feature-stratified round robin over class, GPU count, runtime, submission time, and energy. Experiment 11 exhausts all $4! = 24$ permutations, three penetrations, and two concentration controls. Experiment 18 uses fixed generator-bus electrical-role strata in each public case and the same 0.90 native-load multiplier as the primary preventive DC stress point; its cross-network result is therefore conditional on the declared trace-to-bus scenario, not evidence of utility co-location. Held-out energy-ratio factors span 0.5845–2.6931; the contract uses the capacity-safe upper factor 1.1499 because the raw upper endpoint exceeds the committed 118-MW nameplate. Public Power Grid Library (PGLib) and MATPOWER/PYPOWER cases supply topology, ratings, and costs [15], [32]; RTS-24 follows the published data [33].

B. Baselines, parameters, and inference

Experiment 2 compares eight meter/statistical baselines with the complete-ledger quantile, its exact feasible projection, prevent synthetic control, tail-risk counterfactual, the selected single projection, and risk-constrained verifier. The payment-certified verifier is evaluated separately in Experiments 9 and 15 because its candidate hull and N–1 objective are distinct. An information-set audit verifies feature equality and excludes execution truth. Real-time, elastic, and batch deferral windows are 0, 4, and 512 slots; flexible capacity is 118 MW plus 6 MW fixed demand. Six projection penalties and four risk reserves are validation-only; nested contiguous-fold ordering selects the reserve before locked days are opened, and the four committed-response projection weights are calibrated separately on earlier validation days. The gate contract commits 5 MWh of event-window service and restricts response calibration to positive tariff and projection candidates. Payments are scored by an independent 40-segment evaluator. Moving three-day bootstrap replicates and all 2^{18} block-sign assignments use Holm correction [34], [35]. The sparse solver grows from 5,760 variables/13,424 nonzeros at 96 slots to 36,480/85,488 at 608 slots, with 0.039–0.212-s recorded solve times.

The model-out check uses polar alternating-current optimal power flow (AC OPF) with active/reactive balance, voltage and reactive-generation bounds, and apparent-power ratings [15]. After the intact solve, every credible outage imposes, for each non-reference generator i ,

$$P_{Gi}^{(k)} = P_{Gi}^{(0)}, \quad i \notin \mathcal{G}_{\text{ref}}, \quad k \in \mathcal{C}. \quad (36)$$

The reference generator balances outage-dependent active losses, while $Q_G^{(k)}$ and voltages remain recourse. Apparent-power and voltage limits are checked after each fixed-plan solve, so (36) tests one shared active plan; post-contingency active schedules are not reoptimized. The implementation sets the non-reference bounds to the intact output plus/minus 10^{-8} MW and accepts only a further 10^{-6} -MW solver-residual margin; the largest recorded deviation is 2.52×10^{-8} MW.

Experiments 1–2 test incentives and workload verification; 3–9 test signed settlement, network transfer, risk certification, and N–1 value; 10–12 test AC feasibility, spatial/scale transfer, and continuous-cycle remuneration. Experiments 13–16 provide measured replay, a job-level witness, conversion intervals, and immutable ledger reconciliation. Experiment 17 freezes the event-gate information set, Experiment 18 evaluates a fixed-active-plan AC certificate across four public networks, Experiment 19 solves the indexed job-level counterfactual, Experiment 21 audits fixed-nameplate and capacity-proportional scaling of that primal, and Experiment 22 replays the same indexed witness through N–1 network settlement. Experiment 23 is an independent controlled-event replay of the gate policy, and Experiment 24 is a full finite-outage RTS-24 panel over the frozen profiles. Confirmatory comparisons use the same 54 locked days unless a deterministic network or ledger rule is stated.

TABLE II
LOCKED 54-DAY MECHANISM-ISOLATION RESPONSE RESULTS (DAILY MEANS).

Method	nRMSE	False MWh	Under-credit	F_1
High-5-of-10	4.214	172.709	6.833	0.144
Ridge	2.772	110.666	5.625	0.235
Gradient boosting	2.351	97.057	5.330	0.246
Extra Trees	3.336	132.971	6.947	0.177
Online metadata	2.133	90.589	4.663	0.264
Post-event metadata	2.124	91.603	4.559	0.267
Post-event quantile	0.944	24.156	6.618	0.532
Synthetic control	2.552	101.835	8.427	0.206
Feasible quantile	0.321	2.837	11.555	0.506
Single feasible	0.344	1.849	13.431	0.442
Tail-risk feasible	0.332	2.717	11.940	0.481
Risk-constrained verifier	0.342	1.726	13.482	0.441

V. RESULTS

A. Incentive boundary and independent verification

Experiment 1 solves the entire 8×7 DR-price/call-probability grid. The dual threshold (13) agrees with the independently solved strategic program in all 56 cells. At \$150/MWh and $q = 0.35$, strategic reference scheduling inflates the event baseline by 8.223 MWh. Of 55.452 MWh credited, 44.458 MWh is unsupported by physical reduction, giving a false-credit ratio of 0.802. The profitable region follows the derived marginal boundary and does not use a fitted classifier; the complete grid and its phase diagram are retained in the experiment output.

Experiment 2 evaluates twelve methods on 54 later days; Table II separates the validation tail-risk counterfactual ensemble from the pointwise-safe output. Tail-risk counterfactuals (nRMSE, F_1 , false credit) from the single projection (0.344, 0.442, 1.849) to (0.332, 0.481, 2.717) MWh/day. The final risk-constrained verifier is (0.342, 0.441, 1.726), a 6.68% false-credit reduction with nRMSE change below 0.002. Against the equally informed feasible-quantile projection, false credit falls from 2.837 to 1.726 MWh/day while nRMSE rises from 0.321 to 0.342, locating an explicit risk-accuracy frontier. Source-separated false credit is 7.431/7.580 MWh/day (final/single) for the observed-meter diagnostic and 1.726/1.849 for mechanism isolation; neither is a utility intervention. Total-plus-CVaR, CVaR-only, and the final envelope change (nRMSE, false credit) by $(-0.012, +0.868)$, $(-0.012, +1.946)$, and $(-0.002, -0.124)$ from the single reference. Removing both risk constraints gives 4.936 MWh/day; total-only, CVaR-only, and combined constraints give 2.716, 3.795, and 2.717 MWh/day, respectively.

The independent trace-meter replay is a complementary deployment check: event-window MAE is 2.701/2.849/2.862 MW for feasible quantile/tail-risk counterfactual/risk-constrained verifier on the same 54 locked days (three-day block intervals in Experiment 20). It measures closed-DCGM alignment and remains separate from the mechanism-isolation response panel.

Closest domain baselines are equation-level controls. Spatial-DR, flexibility, receding-horizon, and all-site response give daily (nRMSE, false credit, F_1) = (3.063, 114.573, 0.137), (0.370, 1.549, 0.411), (0.344, 1.849, 0.442), and (0.329, 0, 0.279); temporal-only/joint controls have migration 0/4.328 MWh. All share the ledger, windows, capacities, event slots, and locked days. Recent load-aggregator and production-trace studies are represented only by matched-information analogues, labelled structural rather than faithful software reproductions [12], [13]; the receding-horizon translation follows [11]. Fig. 3 reports distributions and ablations (810 schedules; binding shares 0.310/0.948; slope 1.003).

B. Nodal value and spatial response

Experiment 3 crosses thirteen counterfactuals with five settlement mechanisms. Table III isolates the mechanism with the trace-anchored reference and reports the locked workload rows. Exact signed nodal value has MAE \$151.719 versus \$360.580 for nodal gross response; the corresponding workload errors are \$455.865 and \$421.677. Signed netting reduces the trace error




Fig. 3. Validation-fitted projection weights, constraint-set ablation, specification robustness, and the precision-recall frontier. Every point is produced by a globally solved workload program.

TABLE III
BASE-CASE AND COMPLETE N-1 SETTLEMENT ERRORS (USD/DAY).

Panel	Settlement	Trace MAE	Workload MAE	Trace median/overpay	Workload median/overpay
Base case	Uniform gross	360.528	421.230	—	—
Base case	Nodal gross	360.580	421.677	—	—
Base case	Uniform signed	150.928	454.651	—	—
Base case	Nodal signed linear	152.390	455.842	—	—
Base case	Nodal exact net value	151.719	455.865	—	—
N-1	Base exact	139.221	167.367	79.79/130.42	121.01/107.69
N-1	N-1 linear	2.208	150.461	0.60/1.91	113.43/20.38
N-1	N-1 exact	0.990	150.955	0.87/0.01	113.43/20.18

by \$208.190/day (Holm-adjusted $p = 3.89 \times 10^{-4}$), whereas uniform-to-nodal pricing changes it by $-\$1.463$ ($p = 0.0118$) and linear-to-exact value by $\$0.671$ ($p = 0.8026$). Workload effects are $-\$34.165$, $-\$1.191$, and $-\$0.023/\text{day}$, quantifying the mechanism-specific trade-offs on the same locked workload panel. All 5,616 interval/counterfactual instances satisfy (26); the largest violation is 2.03×10^{-11} dollars.

Experiment 4 applies the predeclared largest gross-minus-net offset (day 87), exposing the rebound component that positive-only local accounting omits.

C. Cross-network and N-1 security evidence

Experiment 5 evaluates 6,480 four-network outcomes; exact value improves the trace-anchored error in all 12 cells (mean $\$1.913/\text{day}$; eight Holm-adjusted tests below 0.05), with a $\$0.050/\text{day}$ workload change. Experiment 6 solves 810 capacity/deadline schedules (binding shares 0.031–0.911; false credit 2.374–2.740 MWh).

Experiment 7 enumerates 16 four-site and 256 eight-site coalitions (maximum budget residual 5.68×10^{-14} dollars) and 1,296 states in the exact 20-participant panel. Experiment 8 enforces (31) for all 37 RTS-24 non-islanding outages: trace/workload MAE is $\$139.221/\$2.208/\$0.990$ and $\$167.367/\$150.461/\$150.955$ for base exact/N-1 linear/N-1 exact, with 1,000 p.u. maximum loading.

Experiment 9 tests (33) with 54 global programs and a validation-frozen seven-profile hull. The contractual cap is selected only by contiguous-fold baseline nRMSE, whereas the payment target is selected by an independent validation-only N-1 payment-MAE criterion; the two roles are recorded separately and neither uses locked outcomes. Held-out factors 0.734/0.996/1.278 have no cap violation above 10^{-8} USD; the independent 40-segment MAE, averaged over these three scenarios, is $\$194.421/\$179.232/\$191.682/\179.589 per day for payment-certified verifier, single, feasible-quantile, and risk-constrained verifier. Experiment 15 gives interval widths $\$11.229/\22.054 at $q_{01}/\text{capacity-safe } q_{99}$ (maxima $\$50.090/\97.951); $q_{99} = 2.6931$ is clipped to 1.150. The unseen-factor transfer panel (0.80/1.20) gives payment-certified MAE 40.759/61.350 and feasible-quantile MAE 38.102/57.350, with no factor used in certification.




Fig. 4. Cross-network preventive AC N-1 certificate. Each point aggregates the predeclared snapshots and all native-case admissible connected outages; the dashed line in (a) is the unit apparent-power limit.

TABLE IV
INFORMATION-BOUNDARY AND LOCKED RESPONSE CERTIFICATES (54-DAY MEANS). THE INDEPENDENT GATE ROW IS GENERATED BY A PREDECLARED EVENT TARIFF AND SERVICE FLOOR ON A MASKED LEDGER; ITS LP DOES NOT RECEIVE THE VERIFIER OUTPUT AND ITS RESPONSE TRAJECTORY IS CHECKED TO DIFFER FROM THE GATE PLAN. IT THEREFORE TESTS CONTRACT/RESPONSE SEPARATION BUT IS NOT A FIELD TREATMENT ESTIMATE. THE COMPLETE-LEDGER ROW IS THE MECHANISM-ISOLATION VERIFIER RESULT.

Method	nRMSE	Gross false MWh	Payable MWh	Meter underpay	Recall	F1
Independent gate-committed response	1.15×10^{-17}	2.63×10^{-16}	224.844	1.97×10^{-16}	1.000	1.000
Complete-ledger risk-constrained verifier (mechanism isolation)	0.342	1.726	10.532	13.482	0.352	0.441

Experiment 10 replaces the linear network with AC OPF [15]: 864 base and 5,400 corrective IEEE-30/39/118 outage solves satisfy the limits; these are the connected, native-case admissible AC cells reported by the panel. Experiment 18 separately evaluates the shared-active-plan certificate on its declared admissible AC set. The 3,888 shared-active-plan cells have zero deviation (maximum loading 0.880/0.907/0.935 p.u. at 3%/6%/9%). Experiment 11 spans 24 regional permutations, three penetrations, four counterfactuals, and 54 days; RiskSafe/single MAE is $\$105.250/\$211.555/\$315.207$ and feasible-quantile is $\$81.500/\$164.024/\$243.866$ per day. Experiment 16 audits 66,769/28,413 training/test jobs (MAE 11.223 W, RMSE 15.818 W, $R^2 = 0.8994$), while Experiment 12 reports complete-cycle/event-window MAE $\$0.641/\$0.658/\$0/\0 and $\$21.882/\$29.287/\$5.325/\5.325 ; all 54 contracts activate.

Fig. 4 summarizes the cross-network preventive certificate: maximum apparent-power loading, voltage-limit deviation, and non-reference active-plan deviation for each penetration and public network.

D. Independent trace and deployment-boundary audits

Immutable MIT scheduler/DCGM replay gives slot MAE 3.740 versus 5.678 MW at 100% coverage and a 71,128-job residual below 10^{-17} MWh. The observational trace-meter MAE is 2.862 MW for risk-constrained verifier (2.701 for feasible quantile); no utility event label is available. Experiment 23 supplies the information-boundary test: a predeclared tariff and an independently solved LP act on a submission-masked ledger without a statistical target, risk envelope, candidate hull, or verifier output. Across 54 locked days the gate-committed response has nRMSE 1.15×10^{-17} , false credit 2.63×10^{-16} MWh/day, payable 224.844 MWh/day, a mean tracking MAE of 0.625 MW, and precision/recall/ $F_1 = 1.000/1.000/1.000$ (Table IV); this is a controlled certificate, not a field treatment estimate. Validation selects $\$10/\text{MWh}$, weight 1.0, and a 20-MWh/day false-credit budget; reserve $\eta = 0.60$ is planning-only. Experiment 18 supplies 9,648 fixed-active-plan AC-OPF outcomes at native-load multiplier 0.90, with zero voltage violation and maximum loading 1.000000 p.u.

E. Indexed workload counterfactual

Experiment 19 assigns a variable to each positive-energy job and admissible slot. Its allocation-runtime-plus-96-slot window, 0.001-MW/GPU bound, native region, exact energy, and 118-MW capacity produce 12,293,445 variables by an exact separable continuous-knapsack decomposition (sparse LP fallback only when a capacity row binds). Native/counterfactual event energy is 0.194/0 MWh, with zero energy residual and




Fig. 5. Independent controlled-event replay over the 54 locked days. The gate-committed response is scored against a separately parameterized, service-floor-constrained tariff LP from the submission-masked ledger; the verifier output is not an input to the event policy.




Fig. 6. Complete frozen-profile RTS-24 security replay. Each cell evaluates all 37 finite non-islanding line outages; the shaded envelope shows the base and worst post-contingency loading across the 54 locked days and eight event slots.

29,492 MWh minimum slack (Table V). Experiment 22 checks the typed invariant before evaluating all 1,056 event slots and every finite RTS-24 N-1 outage; its secure value is \$0.004293 per representative interval and no aggregate reoptimization is used. Experiment 21 reports $3475.108\times$ stress and $1.036\times$ fixed-nameplate scales.

Experiment 23 replays the locked days with an independently generated controlled event: the gate masks post-commitment arrivals and a separate, differently tariffed LP with a predeclared service floor produces the meter. The gate policy has a machine-zero baseline residual, a 0.625-MW mean response-tracking error, and $F_1 = 1.000$ over 54 days, verifying the information boundary without a field causal claim (Fig. 5).

Experiment 24 freezes two profiles and evaluates all 37 finite non-islanding RTS-24 outages in 54 days and 8 event slots (864 cells; 31,968 contingency solves). Every cell succeeds and the maximum post-contingency loading is 1.000000 p.u. (Fig. 6).

TABLE V
EXACT INDEXED JOB-LEVEL COUNTERFACTUAL CERTIFICATE.

Quantity	Value
Positive-energy jobs / job-slot variables	71,128 / 12,293,445
Native / counterfactual event energy (MWh)	0.194 / 0
Gross / net reduction (MWh)	0.194 / 0.194
Event rebound (MWh)	0
Maximum job-energy residual (MWh)	0
Minimum regional-slot capacity slack (MWh)	29,492
Capacity-proportional / fixed-nameplate scale ($\times$)	3475.108 / 1.036
Job-to-network aggregation residual (MWh)	0
Event slots in coupled replay	1,056
Secure N-1 net value (USD/representative interval)	0.004293
Maximum counterfactual post-contingency loading (p.u.)	1.0000

VI. CONCLUSION

This paper couples a provenance-bound workload verifier to signed nodal settlement. The risk-constrained profile reaches nRMSE/false credit 0.342/1.726 MWh/day; a separately solved controlled-event replay has a machine-zero baseline residual, 0.625-MW mean response-tracking error, and $F_1 = 1.000$ without a field treatment claim (Table IV). The 71,128-job witness spans 12,293,445 declared-window variables, yields 0.194 MWh net reduction with zero rebound, and maps to a zero-residual RTS-24 N-1 certificate; the frozen panel succeeds in all 864 cells covering 37 outages. Scale factors $3475.108\times$ and $1.036\times$ distinguish capacity stress from the fixed-nameplate witness. Deployment requires signed co-located telemetry, topology access, and capacity-safe conversion; causal utility validation remains external.

REFERENCES

- [1] California Independent System Operator, "California iso baseline accuracy work group final proposal," tech. rep., California ISO, 2017.
- [2] X. Wang and W. Tang, "Modeling and analysis of baseline manipulation in demand response programs," *IEEE Transactions on Smart Grid*, vol. 13, no. 2, pp. 1178–1186, 2022.
- [3] Z. Liu, A. Wierman, Y. Chen, B. Razon, and N. Chen, "Data center demand response: Avoiding the coincident peak via workload shifting and local generation," *Performance Evaluation*, vol. 70, no. 10, pp. 770–791, 2013.
- [4] M. A. Adnan, R. Sugihara, Y. Ma, and R. Gupta, "Energy efficient geographical load balancing via dynamic deferral of workload," in *2012 IEEE Fifth International Conference on Cloud Computing*, pp. 188–195, 2012.
- [5] Y. Cao, M. Cheng, S. Zhang, H. Mao, P. Wang, C. Li, Y. Feng, and Z. Ding, "Data-driven flexibility assessment for internet data center towards periodic batch workloads," *Applied Energy*, vol. 324, p. 119665, 2022.
- [6] Y. Cao, F. Cao, Y. Wang, J. Wang, L. Wu, and Z. Ding, "Managing data center cluster as non-wire alternative: A case in balancing market," *Applied Energy*, vol. 360, p. 122769, 2024.
- [7] M. T. Takei, M. Qadrdan, J. Summers, and J. Gustafsson, "Data centres as a source of flexibility for power systems," *Energy Reports*, vol. 13, pp. 3661–3671, 2025.
- [8] I. Riepin, T. Brown, and V. Zavala, "Spatio-temporal load shifting for truly clean computing," *Advances in Applied Energy*, vol. 17, p. 100202, 2025.
- [9] W. Zhang, L. A. Roald, A. A. Chien, J. R. Birge, and V. M. Zavala, "Flexibility from networks of data centers: A market clearing formulation with virtual links," *Electric Power Systems Research*, vol. 189, p. 106723, 2020.
- [10] W. Zhang and V. M. Zavala, "Remunerating space-time, load-shifting flexibility from data centers in electricity markets," *Applied Energy*, vol. 326, p. 119930, 2022.
- [11] W. Zhang, L. A. Roald, and V. M. Zavala, "Exploring the impacts of power grid signals on data center operations using a receding-horizon scheduling model," *IEEE Transactions on Sustainable Computing*, vol. 8, no. 2, pp. 245–256, 2023.
- [12] Y. Zhang, A. Tsiligrakis, I. C. Paschalidis, and A. K. Coskun, "Data center and load aggregator coordination towards electricity demand response," *Sustainable Computing: Informatics and Systems*, vol. 42, p. 100957, 2024.
- [13] A. Caprara, Y. Yu, F. Teng, A. Junyent-Ferré, E. Bullich-Massagué, and M. Aragüés-Peñalba, "Data center workload flexibility for power system demand response: Evidence from alibaba traces," *International Journal of Electrical Power and Energy Systems*, vol. 178, p. 111940, 2026.
- [14] B. Stott, J. Jardim, and O. Alsac, "DC power flow revisited," *IEEE Transactions on Power Systems*, vol. 24, no. 3, pp. 1290–1300, 2009.
- [15] R. D. Zimmerman, C. E. Murillo-Sanchez, and R. J. Thomas, "MATPOWER: Steady-state operations, planning, and analysis tools for power systems research and education," *IEEE Transactions on Power Systems*, vol. 26, no. 1, pp. 12–19, 2011.
- [16] D. A. Tejada-Arango, P. Sánchez-Martín, and A. Ramos, "Security constrained unit commitment using line outage distribution factors," *IEEE Transactions on Power Systems*, vol. 33, no. 1, pp. 329–337, 2018.
- [17] National Institute of Standards and Technology and Q. Dang, "Secure hash standard," Tech. Rep. FIPS PUB 180-4, National Institute of Standards and Technology, 2015.
- [18] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge University Press, 2004.
- [19] A. E. Hoerl and R. W. Kennard, "Ridge regression: Biased estimation for nonorthogonal problems," *Technometrics*, vol. 12, no. 1, pp. 55–67, 1970.
- [20] J. H. Friedman, "Greedy function approximation: A gradient boosting machine," *The Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, 2001.
- [21] P. Geurts, D. Ernst, and L. Wehenkel, "Extremely randomized trees," *Machine Learning*, vol. 63, pp. 3–42, 2006.
- [22] A. Abadie, A. Diamond, and J. Hainmueller, "Synthetic control methods for comparative case studies: Estimating the effect of california's tobacco control program," *Journal of the American Statistical Association*, vol. 105, no. 490, pp. 493–505, 2010.
- [23] R. T. Rockafellar and S. Uryasev, "Optimization of conditional value-at-risk," *The Journal of Risk*, vol. 2, no. 3, pp. 21–41, 2000.
- [24] Q. Huangfu and J. A. J. Hall, "Parallelizing the dual revised simplex method," *Mathematical Programming Computation*, vol. 10, no. 1, pp. 119–142, 2018.
- [25] B. Satchidanandan, M. Roohbehi, and M. A. Dahleh, "A two-stage mechanism for demand response markets," *IEEE Control Systems Letters*, vol. 7, pp. 49–54, 2023.
- [26] M. Chen, C. Gao, M. Shahidehpour, and Z. Li, "Incentive-compatible demand response for spatially coupled internet data centers in electricity markets," *IEEE Transactions on Smart Grid*, vol. 12, no. 4, pp. 3056–3069, 2021.
- [27] J. John F. Nash, "The bargaining problem," *Econometrica*, vol. 18, no. 2, pp. 155–162, 1950.
- [28] North American Electric Reliability Corporation, "TPL-001-5.1: Transmission system planning performance requirements," tech. rep., North American Electric Reliability Corporation, 2020.
- [29] S. Samsi, M. L. Weiss, D. Bestor, B. Li, M. Jones, A. Reuther, et al., "The mit supercloud dataset," in *2021 IEEE High Performance Extreme Computing Conference*, pp. 1–8, 2021.
- [30] L. S. Shapley, "A value for n-person games," in *Contributions to the Theory of Games, Volume II* (H. W. Kuhn and A. W. Tucker, eds.), pp. 307–317, Princeton University Press, 1953.
- [31] Y. Wang, Y. Chen, Z. Li, X. Kang, Y. Fang, Y. Zhou, Y. Zheng, Z. Tang, X. He, R. Guo, X. Wang, Q. Wang, A. C. Zhou, and X. Chu, "BurstGPT: A real-world workload dataset to optimize large language model serving systems," in *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pp. 5831–5841, 2025.
- [32] S. Babaeinejadzarokolae, A. Birchfield, R. D. Christie, C. Coffrin, C. DeMarco, R. Diao, et al., "The power grid library for benchmarking ac optimal power flow algorithms," *arXiv preprint arXiv:1908.02788*, 2019.
- [33] C. Grieg, P. Wong, P. Albrecht, R. Allan, M. Bhavaraju, R. Billinton, Q. Chen, C. Fong, S. Haddad, S. Kuruganty, et al., "The IEEE reliability test system-1996," *IEEE Transactions on Power Systems*, vol. 14, no. 3, pp. 1010–1020, 1999.
- [34] H. R. Künsch, "The jackknife and the bootstrap for general stationary observations," *The Annals of Statistics*, vol. 17, no. 3, pp. 1217–1241, 1989.
- [35] S. Holm, "A simple sequentially rejective multiple test procedure," *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979.