

Security-Aware Workload-State Counterfactual Verification and Nodal Net-Value Settlement for Spatially Coupled AI Data Centers

Anonymous Authors

Abstract—Artificial-intelligence (AI) data centers can provide demand response (DR), but a meter-only baseline cannot distinguish a realizable workload shift from strategic deferral, rebound, or a transmission-feasible response. We present a security-aware workload-state verifier and nodal settlement that link a committed job ledger, feasible service trajectories, and contingency-constrained network value. A sparse multi-period linear program enforces release, conservation, deadlines, capacity, migration, and cross-midnight completion. A validation-only convex quadratic program with a linear false-credit/CVaR epigraph and a two-sided event band gives an auditable risk contract; an event-gate mode never credits unsubmitted future work. Signed dispatch-value differences under finite non-islanding single-contingency (N-1) constraints determine payments, with exact coalition accounting for allocation. On 5,188,507 BurstGPT requests and 71,128 positive-energy MIT Supercloud jobs, unsupported credited response falls from 2.026218 to 1.910590 MWh/day (5.71%) and normalized root-mean-square error falls from 0.353737 to 0.348425 (1.50%) over 54 locked days. Exact N-1 settlement reduces trace-anchored error from \$192.69 to \$0.90/day across 37 RTS-24 outages, while end-to-end errors are \$207.67, \$155.25, and \$155.40/day for base-case exact, N-1 linear, and N-1 exact settlement. Independent replay covers every locked 15-minute slot and job-level residuals are below 10^{-17} MWh. Held-out conversion endpoints, four public AC networks, spatial permutations, and a workload-hull payment interval test the declared uncertainty boundary. The non-co-located public traces define a reproducible benchmark, not a utility field event.

Index Terms—AI data centers, counterfactual verification, demand response, security-constrained dispatch, nodal settlement.

I. INTRODUCTION

The growth of large-model training and inference is converting data centers from nearly static industrial loads into large, software-defined grid resources. Interactive inference is latency constrained, elastic inference can be delayed briefly, and batch GPU work can be deferred for hours or migrated between regions. These capabilities enable demand response (DR), but they also invalidate the usual interpretation of a meter decrease. A reduction at one facility can be followed by delayed consumption or simultaneous rebound at another node. Paying the positive local decrease can therefore reward a redistribution that creates little or negative network value.

Conventional DR measurement and verification estimates the unobserved no-event demand from historical meter data. Arithmetic historical baselines remain common in market practice, despite known bias and manipulation incentives [1], [2]. Regression-based facility models quantify baseline uncertainty but do not encode a computing workload ledger [1]. More broadly, DR research establishes the economic role of responsive load, yet a meter-only estimator cannot determine whether queued jobs could actually have produced its predicted trajectory.

Data-center scheduling studies provide the missing physical state. Temporal deferral, dynamic capacity, and geographical

load balancing can reduce energy cost and peak demand [3]. Geographical assignment provides the complementary spatial control mechanism [4]. Recent work assesses data-driven flexibility and aggregator coordination [5], [6], while multi-level dispatch refines delay-sensitive electrical response. Hierarchical optimal-consumption models further connect facility efficiency with response quality [7]. Virtual-link market clearing and space-time remuneration establish how geographical and temporal transfers can enter electricity markets [8], [9]. Receding-horizon operation and uncertainty-aware frequency-regulation coordination further establish the need to retain future arrivals and multi-site state [10]. However, these methods optimize an operator-controlled schedule rather than verify an independently submitted DR counterfactual against an auditable workload ledger. This separation between computing feasibility and market verification motivates the auditable ledger used here.

The second gap is electrical. Geographic workload control is commonly valued with energy prices or gross megawatts. Such accounting ignores the fact that nodal marginal value changes when congestion or a contingency constraint binds. Direct-current (DC) power flow and linear distribution factors provide transparent network sensitivities [11], [12]; N-1 dispatch can impose every credible post-contingency line limit through line-outage distribution factors (LODFs) [13]. No existing data-center DR verification framework jointly certifies workload feasibility, preserves signed spatial response, and settles the resulting security-aware network value. Public workload and facility telemetry can support an observational replay of that ledger, but it does not constitute a causal utility event; a rigorous study must therefore separate measured-execution replay, counterfactual construction, and network-value verification rather than present them as one field intervention.

This paper makes three power-system contributions. 1) We formulate a provenance-bound workload-state demand model in which release, deadline, conservation, cross-midnight completion, capacity, and migration constraints are enforced simultaneously. The resulting job-level witness links each joined job to its immutable scheduler/DCGM record and preserves measured execution energy. 2) We construct a validation-only, continuous-arrival counterfactual verifier with a pointwise credit envelope. The output is an attainable spatial-temporal load trajectory rather than an unconstrained meter-only prediction, and its statistical performance is evaluated under an explicit false-credit/under-credit tradeoff. 3) We develop signed space-time nodal remuneration and a bilateral net-value contract under preventive N-1 dispatch, and map the feasible workload set into a payment uncertainty interval over held-out workload-to-power conversion. The novelty is the coupling invariant that carries one committed workload state into physical verification, network valuation, and a set-valued settlement rule; the individual LP, dispatch, and allocation

primitives are standard and are not claimed as new in isolation. The certificate is conditional on trusted scheduler and telemetry provenance: a cryptographic digest verifies record integrity after commitment but does not authenticate semantically false fields.

The remainder of the paper is organized as follows. Section II defines the spatial workload and network settlement problem. Section III presents the verifier, security-aware value rule, and allocation mechanism. Section IV specifies the data, baselines, and locked protocol. Section V reports the core counterfactual, execution-replay, job-level, security, payment, continuous-horizon, and provenance experiments. Section VI concludes with the deployment boundary and future validation requirements.

II. SYSTEM MODEL AND PROBLEM FORMULATION

A. Spatial workload state

Let $\mathcal{S}$, $\mathcal{K}$, $\mathcal{D}$, and $\mathcal{T}$ denote source regions, workload classes, destination data centers, and 15-minute intervals. The measured ledger supplies arrival energy a_{skt} for source s , class k , and release interval t . The decision $x_{skd\tau} \geq 0$ is the energy served at destination d in interval τ . Its cumulative state is

$$y_{skt} = \sum_{\tau=0}^t \sum_{d \in \mathcal{D}} x_{skd\tau}. \quad (1)$$

With cumulative arrivals $A_{skt} = \sum_{\tau=0}^t a_{sk\tau}$ and deadline D_k , release and deadline feasibility are

$$y_{skt} \leq A_{skt}, \quad (2)$$

$$y_{skt} \geq A_{sk,t-D_k}, \quad t \geq D_k. \quad (3)$$

The state recursion equals service, terminal service equals total arrivals, and $\sum_{s,k} x_{skdt} \leq \bar{P}_d \Delta t$. Facility demand is

$$p_{dt} = P_d^{\text{fix}} + \Delta t^{-1} \sum_{s,k} x_{skdt}. \quad (4)$$

For rolling window end T^+ and final event-day release T^0 , the terminal state is not forced to consume unexpired future work. Instead,

$$y_{skT^+} = A_{sk,\max\{T^0, T^+ - D_k\}}, \quad (5)$$

which completes every job due inside the window and every event-day arrival, while retaining later real arrivals in the release constraints. The window starts $D_{\max} + 96$ intervals before the event day and ends $D_{\max}$ intervals after it; omitted earlier work must therefore expire before the evaluated day. Equations (2) and (3) prevent early execution and late completion, while (4) maps the feasible service state to the facility meter. This model follows workload deferral and geographical assignment foundations [3], [4]; its rolling implementation follows the state-retention principle in [10]. Waiting and migration costs are $w_k(\tau - t)x_{skd\tau}$ and $m_{sd}x_{skd\tau}$, respectively. Because arrivals are fixed, replacing waiting duration by service time adds only a constant.

For provenance, each retained job is represented by the canonical tuple and its digest in

$$\begin{aligned} \ell_j &= (\text{id}_j, t_j^{\text{submit}}, t_j^{\text{start}}, t_j^{\text{end}}, E_j, n_j^{\text{tel}}, q_j, \kappa_j, \sigma_j), \\ \mathcal{L} &= \text{sort}\{\ell_j : j \in \mathcal{J}\}, \\ h_{\mathcal{L}} &= \text{SHA256}(\text{serialize}(\mathcal{L})). \end{aligned} \quad (6)$$

where E_j is the positive aggregated Data Center GPU Manager (DCGM) energy and n_j^{tel} is its raw telemetry-record count. The digest uses the Secure Hash Standard [14]; the canonical row order, one-to-one join, release-execution order, and raw-to-join energy equality are checked independently of the scheduling linear program (LP). Consequently, the flow witness certifies the exact committed ledger and measured profile, while the

digest does not make an unverified claim about jobs that might have been submitted before the commitment. Experiment 16 reports this provenance certificate and reconciles the measured regional capacity envelope with the 118-MW nameplate that was committed before the validation/test split; the envelope is not used to select that nameplate.

The trust boundary is explicit. The verifier assumes that the scheduler and DCGM services are the authoritative sources for the fields used in ℓ_j , and that the digest is committed before the event decision. Under this assumption, a changed record is detectable and a joined record cannot be silently replaced after commitment. The digest does not, however, establish that a semantically false deadline, region label, service class, or energy value was truthful when it was signed. Authenticity of those fields requires an upstream scheduler attestation, meter signature, or trusted execution interface; it is outside the optimization certificate proved here.

B. Counterfactual, response, and false credit

For event set $\mathcal{E} \subset \mathcal{T}$, p_{dt}^0 denotes the no-event counterfactual and p_{dt}^1 denotes the event trajectory generated by the declared workload optimization. The mechanism-isolation experiments use $p_{dt}^{1,\text{sim}}$ for this optimized response, whereas Experiment 13 uses $p_{dt}^{1,\text{meas}}$ for the independently measured MIT/DCGM execution trace. We reserve p_{dt}^{obs} for measured execution used only as scoring truth; it is never supplied to an estimator or used to claim a causal utility event. For either response trace, $r_{dt} = p_{dt}^0 - p_{dt}^1$ is signed: gross local response is $[r_{dt}]_+$, while net response retains negative rebound. For an estimate $\hat{p}^0$, unsupported credited energy is

$$F = \Delta t \sum_{d,t \in \mathcal{E}} [\hat{p}_{dt}^0 - \max\{p_{dt}^{\text{obs}}, p_{dt}^1\}]_+. \quad (7)$$

The audit computes (7) separately for every locked day so that a favorable aggregate error cannot hide unsupported positive credit. We also report normalized root-mean-square error (nRMSE), bias, precision, recall, F_1 , under-credit, payment error, and overpayment.

C. Strategic reference scheduling

Consider an arithmetic mean of $H = 10$ reference days, event probability q , and DR price π^{DR} . If reference day j adds event-window service δR_j , the expected payment change is

$$\Delta \Pi = \frac{q\pi^{\text{DR}}}{H} \sum_{j=1}^H \delta R_j - \sum_{j=1}^H \Delta C_j. \quad (8)$$

Let c'_j be the right derivative of the minimum operating cost with respect to an event-window minimum-service constraint on day j . Linear-program sensitivity [15] gives the following exact local condition:

$$\frac{q\pi^{\text{DR}}}{H} > \min_j c'_j \iff \exists \text{ a strictly profitable local deviation.} \quad (9)$$

The historical-baseline premise is documented in [1]; (8)–(9) are this paper's workload-specific specialization.

D. Network value

Let M be the data-center-to-bus incidence matrix and $p_t = p_t^{\text{native}} + Mp_t^{\text{dc}}$ the nodal demand vector. For generator-to-bus map C_g , generator bounds, power transfer distribution factor

(PTDF) H , and line ratings $\bar{f}$, the piecewise-linear security-constrained economic dispatch (SCED) value is

$$V(p) = \min_g \sum_i C_i(g_i) \quad (10)$$

$$\text{s.t. } \mathbf{1}^\top g = \mathbf{1}^\top p, \quad \underline{g} \leq g \leq \bar{g},$$

$$-\bar{f} \leq H(C_g g - p) \leq \bar{f}.$$

The incidence mapping aggregates device demand at its connection bus, and (10) is the standard lossless DC representation [11], [12]. The realized signed grid value is

$$\Delta V = V(p^0) - V(p^1). \quad (11)$$

Positive ΔV is a credit and negative ΔV a debit; no positive-only truncation is applied.

III. PROPOSED METHODOLOGY

A. Framework and matched-information pre-estimation

Fig. 1 shows the complete chain. Meter history and the submitted-job ledger enter a statistical pre-estimator and the workload feasible set in parallel. The pre-estimator never receives execution energy used as scoring truth. For each event day, it uses only the information available under one of two declared protocols: lagged meter/calendar features for online baselines, or the complete submitted ledger for equally informed post-event comparisons. Ridge regression and gradient boosting follow [16], [17]; extremely randomized trees follow [18]. Two stronger comparators are also solved: median-loss gradient boosting with the complete submitted ledger and a nonnegative synthetic control whose weights globally minimize pre-event absolute error over 28 donor days [19]. To separate the effect of projection from that of pre-estimation, the complete-ledger quantile output is also passed through the same exact workload-feasible projection, with its penalty chosen by contiguous validation folds. The same complete ledger is supplied to the post-event learner, the selected single projection, and the proposed convex verifier.

B. Sparse cumulative workload-state verifier

For statistical estimate $\tilde{p}$ and each predeclared penalty ρ_ℓ , the first-stage program solves

$$\min_{x,y,z} \mathcal{C}_w(x) + \rho_\ell \sum_{d,t} z_{dt} \quad \text{s.t.} \quad z_{dt} \geq \pm(p_{dt}(x) - \tilde{p}_{dt}), \quad (12)$$

subject to (1)–(4). Absolute-value epigraphs are standard convex constructions [15]; using them to project a meter baseline onto an auditable workload ledger is specific to this work. Here $\mathcal{C}_w$ denotes workload operating cost; C_i denotes generator-segment cost in (10), and $V(p)$ denotes the network dispatch value. Thus (12) changes the objective, not the physical feasible set. Six globally solved schedules $p^{(\ell)}$ share identical arrivals and constraints. Their validation-only convex combination $\bar{p} = \sum_\ell \alpha_\ell p^{(\ell)}$, $\alpha \geq 0$, $\mathbf{1}^\top \alpha = 1$, remains feasible because the common scheduling set is convex.

The coefficients minimize validation squared error subject to both total and daily-tail false-credit budgets, where conditional value at risk (CVaR) measures the latter [20]. Let ℓ^* be the single candidate with minimum worst contiguous-fold nRMSE and B_j^{ref} its validation false exposure on day j . For reserve $\beta \in (0, 1]$,

$$\sum_j F_j(\alpha) \leq \beta \sum_j B_j^{\text{ref}}, \quad (13)$$

$$\text{CVaR}_{0.75}(F_j(\alpha)) \leq \beta \text{CVaR}_{0.75}(B_j^{\text{ref}}). \quad (14)$$

For implementation, the positive part is not differentiated through a sort. For each event sample n and validation day j , the program introduces $f_n \geq 0$, a CVaR threshold $\nu \geq 0$, and tail slacks $u_j \geq 0$:

$$f_n \geq p_n^\alpha - \max\{p_n^0, p_n^1\}, \quad F_j = \sum_{n \in j} f_n, \quad (15)$$

$$F_j - \nu - u_j \leq 0, \quad (16)$$

$$\nu + \frac{1}{0.25J} \sum_j u_j \leq \beta \text{CVaR}_{0.75}(B_j^{\text{ref}}). \quad (17)$$

Together with the simplex equality, box bounds on α , and the total budget in (13), these are linear constraints and the squared error objective is a convex quadratic program. A HiGHS feasibility solve is followed by a sparse trust-region solve; the reported KKT and primal residuals are below 10^{-8} in every fitted fold. Four contiguous folds select β using validation nRMSE, then risk ratio; held-out noninferiority is reported as an outcome and is not used as a hidden constraint. A second exact projection targets $\bar{p}$ while imposing, for every event sample,

$$p_{dt}^{\text{safe}} \leq p_{dt}^{(\ell^*)}. \quad (18)$$

a) *Two-sided credit band.*: The upper envelope alone would admit an arbitrarily conservative zero-credit profile. We therefore impose the matching lower band on the event slots,

$$p_{dt}^{(\ell^*)} - \varepsilon \leq p_{dt}^{\text{safe}} \leq p_{dt}^{(\ell^*)}, \quad (d, t) \in \mathcal{D} \times \mathcal{E}, \quad (19)$$

where $\varepsilon = 1$ MW is fixed before the locked test split. The lower bound is implemented together with the workload constraints, not as a post-processing clip. Thus the output remains an attainable schedule while its credited event trajectory is within a declared band of the selected single projection. The total constraint (13) controls aggregate exposure and (14) controls the worst daily tail; the pair (18)–(19) supplies the samplewise credit certificate. Since $[z - c]_+$ is monotone, the upper inequality bounds false credit for any unseen actual trajectory, while the lower inequality bounds the additional under-credit by at most $\varepsilon |\mathcal{D}| |\mathcal{E}| \Delta t$ relative to the selected reference. The selected single feasible projection proves nonemptiness, and the locked results report both margins explicitly.

b) *Decision-time information boundary.*: The ex-post selector can use the complete submitted ledger because it is an offline settlement protocol. For an event-gate deployment, all arrivals after the gate are removed before optimization and terminal completion is required only for work due by the gate. We additionally solve a committed-ledger safe mode with $p_{dt} \leq P_d^{\text{fix}}$ on the event slots. This mode is an exact workload LP, uses no future arrivals, and credits only service backed by the committed ledger; later arrivals become eligible only after commitment. It is reported separately from the information-boundary diagnostic and the ex-post benchmark.

The cumulative state reduces the constraint-matrix nonzeros from $O(|\mathcal{S}| |\mathcal{K}| |\mathcal{D}| |\mathcal{T}|^2)$ for explicit prefixes to $O(|\mathcal{S}| |\mathcal{K}| |\mathcal{D}| |\mathcal{T}|)$. A 512-slot deadline is handled by (5) with every observed future arrival retained; no zero-arrival completion buffer is used in the final continuous-horizon certificate. The day trajectory is embedded by a lexicographic pair of linear programs: minimum L1 distance first, followed by minimum operating cost on that optimal-distance face. Hence no penalty weight trades physical fit against cost. All linear programs are solved to global optimality with HiGHS [21]. The resulting state-to-credit sequence is summarized in Fig. 2; the execution replay and interval audit are reported separately in Experiments 13–15.




Fig. 1. Verification-to-settlement framework. Validation labels select the counterfactual risk contract; locked execution truth is used only for scoring. The right branch computes aggregate N-1 value and exact participant shares.




Fig. 2. Sparse cumulative workload-state verifier. Real release-to-deadline edges feed the global feasibility LP; the pointwise credit envelope and the finite/interval payment gates are applied only after the workload state is certified.

C. Signed nodal and N-1 settlement

Let $\lambda^0 \in \partial V(p^0)$ be the nodal demand-value subgradient. Signed linear settlement is $(\lambda^0)^\top (p^0 - p^1)$, while (11) is the realized system-value benchmark and the payoff of an explicit bilateral avoided-cost contract; it is not relabeled as an independent system operator (ISO) market payment. The convex subgradient inequality [15] yields the global certificate $0 \leq (\lambda^0)^\top (p^0 - p^1) - \Delta V = V(p^1) - V(p^0) - (\lambda^0)^\top (p^1 - p^0)$. (20)

Thus the linear rule globally upper-bounds exact value in the implemented polyhedral market model. Equality holds only when both endpoints share a supporting dual subgradient.

For market-compatible remuneration, let $\mathcal{W}$ contain the complete response cycle from the earliest deadline-feasible anticipatory shift through the latest recovery interval. The space-time rule is

$$\Pi^{\text{ST}} = \Delta t \sum_{t \in \mathcal{W}} \sum_d \lambda_{dt}^0 (p_{dt}^0 - p_{dt}^1). \quad (21)$$

This is the nodal virtual-link remuneration structure of [8], [9], applied here to a verified counterfactual. Summing site payments reproduces (21) exactly.

Space-time energy remuneration is only one component of the contracted DR product. Let $\mathcal{D}^{\text{part}}$ be the predeclared participating sites, $\mathcal{E}$ the event intervals, and r^{DR} the cleared service price. The signed delivered capacity service and total operator value are

$$Q^{\text{cap}} = \Delta t \sum_{d \in \mathcal{D}^{\text{part}}} \sum_{t \in \mathcal{E}} (p_{dt}^0 - p_{dt}^1), \quad (22)$$

$$G = r^{\text{DR}} Q^{\text{cap}} + \Pi^{\text{ST}}, \quad (23)$$

where p^0 is the minimum-cost no-event schedule under the identical continuous arrival window. Capacity credits and ex-post participation constraints follow established DR mechanism design [22], [23]; the signed Π^{ST} term separately debits spatial transfer and recovery. Equations (22)–(23) jointly define the signed service and operator value.

Let x^1 and x^0 denote the workload service schedules inducing p^1 and p^0 , respectively. Define $C_{\text{opp}} = \mathcal{C}_w(x^1) - \mathcal{C}_w(x^0)$ as the participant opportunity cost and $S = G - C_{\text{opp}}$ as the transferable transaction surplus. With disagreement utilities equal to zero, the symmetric Nash solution [24] activates the

contract iff $S \geq 0$ and sets

$$z = \mathbb{I}\{S \geq 0\}, \quad P^B = z \left(C_{\text{opp}} + \frac{S}{2} \right), \quad u_{\text{dc}} = u_{\text{op}} = z \frac{S}{2}. \quad (24)$$

Thus voluntary nonactivation is the outside option rather than an omitted outcome. For every activated contract, both utilities are nonnegative and the single transfer P^B is exactly budget balanced. The bilateral transfer is reported separately from both ISO nodal remuneration and the exact dispatch value benchmark. Unlike event-only accounting, $\mathcal{W}$ debits both pre-event load increases and delayed rebound; workload conservation gives $\sum_{d,t \in \mathcal{W}} (p_{dt}^0 - p_{dt}^1) \Delta t = 0$ up to solver tolerance.

Security-aware settlement augments (10) with every finite non-islanding single-line outage. If $L_{\ell k}$ is the LODF from outage k to monitored line ℓ , then

$$-\bar{f}_\ell \leq (H_\ell + L_{\ell k} H_k)(C_g g - p) \leq \bar{f}_\ell, \quad \ell \neq k. \quad (25)$$

Equation (25) is the standard preventive N-1 LODF formulation [13]; the experiment enforces all credible rows simultaneously, rather than screening contingencies. Native public ratings are retained. Islanding outages are reported and excluded because a connected single-reference PTFD cannot represent separate islands, consistent with the scope of DC security models and the N-1 planning criterion [25].

D. Scenario-robust N-1 payment certificate

The pointwise envelope controls credited energy but does not by itself certify network payment because spatially different demand vectors can activate different security constraints. Let $\mathcal{P} = \text{conv}\{p^{(1)}, \dots, p^{(L)}\}$ be the convex hull of complete workload-feasible schedules and let $p^{\text{ref}} \in \mathcal{P}$ be the validation-selected single feasible projection (the selected $\rho = 1$ member of the first-stage workload hull). The feasible-quantile projection is retained as an external matched transfer comparator in Section V. Let Ξ_3 contain the 10th, 50th, and 90th percentiles of held-out per-job measured-to-predicted GPU energy. These predeclared empirical scenarios define the finite power-conversion contract supported by independent telemetry [26]. For each settlement day, the payment-certified verifier defines $B_\xi^{\text{ref}} = \Delta t \sum_{t \in \mathcal{E}} V_{N-1}(p_t^{\text{native}} + M\xi p_t^{\text{ref}})$ and solves

$$\begin{aligned} \text{lexmin}_{\alpha, e, \eta, \{g_t^\xi\}} \quad & \left(\sum_{d,t \in \mathcal{E}} e_{dt}, -\eta \right) \\ \text{s.t.} \quad & p^\alpha = \sum_{\ell=1}^L \alpha_\ell p^{(\ell)}, \\ & \alpha \geq 0, \quad \mathbf{1}^\top \alpha = 1, \quad 0 \leq \eta \leq 1, \\ & e_{dt} \geq \pm(p_{dt}^\alpha - p_{dt}^{\text{safe}}), \\ & g_t^\xi \text{ satisfies (10) and (25),} \\ & \quad \text{with } p_t = p_t^{\text{native}} + M\xi p_t^\alpha, \\ & \quad \quad t \in \mathcal{E}, \quad \xi \in \Xi_3, \\ & \Delta t \sum_{t \in \mathcal{E}} \sum_i C_i(g_{it}^\xi) \\ & \quad \leq (1 - \eta) B_\xi^{\text{ref}}, \quad \xi \in \Xi_3. \end{aligned} \quad (26)$$

Generator costs use the same declared piecewise-linear epigraph as market settlement. Two global linear programs implement the lexicographic objective: the first obtains the minimum trajectory error, and the second preserves that optimum within solver tolerance while maximizing the common fractional margin η . Existence of the primal dispatches in (26) is equivalent to

the corresponding upper bound on the optimal N-1 value; the unit weight on p^{ref} proves nonemptiness in every scenario. Consequently, for any $\xi \in \Xi_3$ and realized event load $p^{1,\xi}$, daily payment obeys

$$\begin{aligned} & \sum_{t \in \mathcal{E}} \Delta t [V_{N-1}(p_t^{\text{native}} + M\xi p_t^\alpha) - V_{N-1}(p_t^{1,\xi})] \\ & \leq \sum_{t \in \mathcal{E}} \Delta t [V_{N-1}(p_t^{\text{native}} + M\xi p_t^{\text{ref}}) - V_{N-1}(p_t^{1,\xi})]. \end{aligned} \quad (27)$$

The common realized-cost term cancels, so the certificate requires no execution truth and holds samplewise. Both programs contain every finite non-islanding contingency. The formulation follows the standard SCED primal and linear-program epigraph foundations [12], [15]. Here $\mathcal{P}$ contains the six independently solved first-stage projections from (12) and the matched feasible-quantile projection. The risk-constrained verifier is an external target in the absolute-deviation objective and is not a selectable candidate; the quantile schedule is used only as an independent transfer comparator, not as the contractual payment cap.

The finite cap is complemented by a set-valued payment certificate. For a realized counterfactual value, define $\Delta V_\xi(p) = V_{N-1}(p, \xi) - V_{N-1}(p^{1,\xi})$ and let $\mathcal{P}_{\text{seg}} = \{p^{\text{ref}}, p^{\text{cert}}\}$ denote the two validation-frozen endpoints of the contractual segment. The induced interval is

$$\mathcal{I}_\xi = \left[\min_{p \in \mathcal{P}_{\text{seg}}} \Delta V_\xi(p), \max_{p \in \mathcal{P}_{\text{seg}}} \Delta V_\xi(p) \right]. \quad (28)$$

The selected reference payment is an element of $\mathcal{I}_\xi$ by construction. The interval endpoints use only these two feasible segment endpoints and the realized counterfactual needed for settlement; no oracle baseline is used to select or certify the interval. Thus (28) makes uncertainty in the workload-to-power map visible to the contract instead of relabeling a single forecast as true value. It is not a lower-error claim: interval width is a reported uncertainty cost; the larger seven-profile hull remains the finite-scenario payment certificate in (26).

E. Participant allocation

For participants N , define $v(S)$ as signed dispatch value when coalition S moves from baseline to actual demand and nonmembers remain at baseline. The exact allocation is

$$\phi_i(v) = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(|N| - |S| - 1)!}{|N|!} [v(S \cup \{i\}) - v(S)]. \quad (29)$$

This is the Shapley value [27]; the characteristic function based on signed grid value is the proposed specialization. Grouping (29) by coalition size cancels every intermediate value and gives $\sum_i \phi_i = v(N)$ when $v(\emptyset) = 0$. We enumerate all 16 coalitions for four sites and all 256 coalitions for eight contractual portfolios; no permutation estimator is used. For larger auditable contracts, exchangeable slices at the same electrical site are grouped by their member count. The exact sum retains the binomial multiplicity of every labeled coalition and evaluates only $\prod_g (m_g + 1)$ distinct count states. This symmetry reduction remains the Shapley value and is evaluated through 20 participants.

F. Analytical guarantees

Proposition 1 (manipulation boundary): Under a ten-day arithmetic baseline, differentiable one-sided optimal costs, and an interior feasible increase in at least one reference-day event service, a strictly profitable local manipulation exists if and

only if (9) holds. To prove the statement, differentiate (8) along any feasible reference-day direction. The marginal payment is the same $q\pi^{\text{DR}}/H$ for every day, whereas the least physical cost is $\min_j c'_j$. A strict excess selects a positive direction on an argmin day; if no strict excess exists, every nonnegative local direction has nonpositive first-order profit. Equality admits indifference but does not imply strict profit. Experiment 1 computes c'_j from the honest program's right-hand-side dual and solves the strategic program independently, preventing the behavioral result from being reused as its own certificate.

Proposition 2 (workload and two-sided credit feasibility): For fixed arrivals, every convex combination of solutions to (1)–(4) is feasible, and the envelope output satisfies the deterministic two-sided band (19) on every event cell. Consequently, for every possible unseen execution c_{dt} , its false-credit energy is no larger than that of $p^{(\ell^*)}$, while its additional under-credit is bounded by $\varepsilon|\mathcal{D}||\mathcal{E}|\Delta t$ relative to that reference. Affine equalities are preserved by the simplex coefficients; each linear inequality is preserved by convexity. The upper inequality in (19) and monotonicity of $[z - c]_+$ give

$$[p_{dt}^{\text{safe}} - c_{dt}]_+ \leq [p_{dt}^{(\ell^*)} - c_{dt}]_+.$$

The lower inequality gives $[c_{dt} - p_{dt}^{\text{safe}}]_+ \leq [c_{dt} - p_{dt}^{(\ell^*)}]_+ + \varepsilon$. Summation over sites and event intervals proves both bounds. The result is a feasibility-and-band certificate determined by the declared constraints and ε , not by the distribution of the locked test set.

Proposition 3 (security-aware value consistency): For the piecewise-linear base-case or N–1 dispatch, exact signed settlement equals the corresponding optimal-value difference and the nodal linear rule upper-bounds it by (20). The feasible-set right-hand side is affine in nodal demand, so the optimal value is convex and piecewise affine. Applying the subgradient inequality at p^0 gives (20); applying it again at p^1 yields

$$\begin{aligned} 0 \leq \Delta V_{\text{lin}} - \Delta V &\leq (\lambda^1 - \lambda^0)^\top (p^1 - p^0) \\ &\leq \|\lambda^1 - \lambda^0\|_* \|p^1 - p^0\|. \end{aligned}$$

The same proof applies after stacking all rows in (25), because they remain linear. Gross positive-only settlement has no analogous identity: truncation changes the response vector before it enters the value function.

Proposition 4 (robust daily payment noninferiority): Every solution of (26) is workload feasible and its exact N–1 payment is no larger than the reference payment for every realized event trajectory and every declared $\xi \in \Xi_3$. Workload feasibility follows from Proposition 2. For each ξ , the embedded dispatch is a feasible witness whose cost satisfies the corresponding last constraint in (26); optimal N–1 dispatch can only have lower cost. Subtracting the same scenario-specific realized cost gives (27). Conditional on the telemetry-derived finite set Ξ_3 , the proof is deterministic and contains no test-outcome assumption. *Proposition 5 (continuous worst-case interval certificate):* For a fixed workload profile p , the optimal piecewise-linear N–1 dispatch value $V_{N-1}(p; \xi)$ is convex in a scalar conversion factor $\xi \in [\xi_L, \xi_U]$. Therefore

$$\max_{\xi \in [\xi_L, \xi_U]} V_{N-1}(p; \xi) = \max\{V_{N-1}(p; \xi_L), V_{N-1}(p; \xi_U)\}.$$

The value function is the supremum of affine dual objectives in ξ , hence convex; a convex function on a compact interval attains its maximum at an endpoint. If the two endpoint inequalities $V_{N-1}(p^\alpha; \xi_e) \leq V_{N-1}(p^{\text{ref}}; \xi_e)$ hold for $e \in \{L, U\}$, then the certified candidate also has no larger worst-case interval cost than the reference. This interval result is a worst-case bound; pointwise dominance at every interior ξ is asserted only for the finite set Ξ_3 in Proposition 4. Experiment 15 evaluates

TABLE I
COMPLETE DATA AND EVALUATION PROTOCOL.

Item	Declared use and count
BurstGPT	5,188,507 requests; 2.176 billion tokens; inference arrivals and observed service
MIT scheduler	287,069 rows; submission time, duration, GPU count, and queue state
MIT GPU telemetry	96,893 rows; 68,664 tensor-window jobs and 71,128 full-horizon positive-energy jobs; measured execution energy
Time resolution	15 min; 121 calendar days; 121 valid complete days
Evaluation	16 validation days followed by 54 untouched test days
Networks	RTS-24, IEEE-39, Power Grid Library (PGLib)-118, and IEEE-300 with native ratings
Provenance	Immutable scheduler/DCGM join digest, raw-to-join energy conservation, and measured regional capacity envelope

both endpoints with an independent 40-segment N–1 solver and reports this distinction explicitly.

Proposition 6 (workload-segment payment interval): For a fixed conversion factor and realized counterfactual, the endpoints in (28) are the exact minimum and maximum payments over the declared two-endpoint feasible segment. This follows directly from the finite-set minimum and maximum; it is an exact uncertainty statement for that contractual segment, not an assertion that the segment contains every physically possible schedule. Experiment 15 reports endpoint width for both held-out conversion endpoints; no oracle baseline is used to select the segment.

Proposition 7 (provenance-bound workload witness): Suppose the canonical tuple in (6) is computed from the retained scheduler–DCGM join after the declared duplicate and positive-energy rules. If the certificate verifies one-to-one joining, release–execution ordering, raw-to-join energy conservation, and zero LP residuals, then the job-level flow witness is a feasible representation of that committed ledger. Under the collision-resistance assumption of SHA-256, any post-commit modification of a tuple or its canonical ordering changes the digest except with negligible probability. The statement binds the computational witness to the committed records; it does not certify that an uncommitted pre-event submission was truthful. Experiment 16 checks each premise independently and records the raw source hashes, so the provenance claim is separate from the scheduling solver's own feasibility status.

IV. EXPERIMENTAL DESIGN

A. Data, preprocessing, and information boundary

Table I reports the complete data inventory. BurstGPT contributes 5,188,507 successful requests and observed token service [28]. MIT Supercloud contributes scheduler submissions and independent Data Center GPU Manager energy. The complete positive-energy scheduler–DCGM join contains 71,128 jobs, of which 68,664 start within the common 121-day tensor window used by Experiments 1–13 [26]. Submission time creates queue arrivals; measured execution intervals create scoring truth. A held-out GPU-power regression uses 28,413 observations and obtains $R^2 = 0.899$, mean absolute error (MAE) 11.22 W, and root-mean-square error (RMSE) 15.82 W. Per-job measured-to-predicted energy ratios define Ξ_3 in (26); no synthetic row replaces an invalid record. The 68,664 count is the 121-day workload tensor used by Experiments 1–13. Experiment 14 does not truncate at that tensor boundary; it extends the deadline horizon to the last observed scheduler completion and retains all 71,128 positive-energy joined jobs for its independent job-level flow witness.

The public traces do not contain utility bus locations. Consequently, the four-region coupling is a trace-driven benchmark with fixed, predeclared bus placements and a declared workload-to-power conversion, not a utility field trial. Immutable row identifiers determine spatial partitions before any outcome is computed. Held-out measured-to-predicted job-energy ratios span 0.5845–2.6931 (1st–99th percentiles); Experiments 9, 11, and 15 evaluate this interval rather than treating the 118-MW scale as a measured geographic fact. This design permits exact repeated mechanism comparisons but does not identify a particular operator’s geography. Public PGLib and MATPOWER/PYPOWER cases provide topology, capacities, costs, and ratings [12], [29]. The RTS-24 benchmark follows the published system data [30].

B. Baselines, parameters, and inference

Experiment 2 compares High-5-of-10, ridge regression, gradient boosting, Extra Trees, online metadata gradient boosting, equally informed post-event metadata gradient boosting, complete-ledger quantile boosting, pre-event synthetic control, the exact feasible projection of the complete-ledger quantile, the tail-risk counterfactual, the independently selected single feasible projection, and the risk-constrained verifier. The first eight are meter/statistical comparators; the remaining methods receive the declared ledger as appropriate. The payment-certified verifiers are evaluated separately in Experiments 9 and 15 because its N–1 payment certificate uses a different objective and a frozen candidate hull. An information-set audit verifies feature equality and confirms that no method observes execution truth.

Real-time inference has zero deferral, elastic inference four slots, and batch GPU work 512 slots. Per-site flexible capacity is 118 MW plus 6 MW fixed facility demand; the 118-MW nameplate is committed before the validation/test split and Experiment 16 only reconciles it with the complete observed MIT/DCGM batch envelope. Six projection penalties $\{0.03, 0.1, 0.3, 1, 3, 10\}$ and four risk reserves $\{0.60, 0.75, 0.90, 1.00\}$ are selected only on validation data. The settlement uses 10 generator-cost segments and is scored by an independently solved 40- or 80-segment evaluator. All 5,000 bootstrap replicates use moving three-day blocks [31]. Paired tests enumerate all 2^{18} signs of 18 nonoverlapping three-day blocks; Holm correction controls family-wise error [32].

The nonlinear model-out check uses the standard polar AC optimal power flow with active/reactive balance, voltage bounds, generator reactive bounds, and apparent-power ratings [12]. After solving the intact state, the preventive panel imposes, for every non-reference generator i and every credible outage k ,

$$P_{Gi}^{(k)} = P_{Gi}^{(0)}, \quad i \notin \mathcal{G}_{\text{ref}}, \quad k \in \mathcal{C}. \quad (30)$$

The reference generator balances only outage-dependent active losses; $Q_G^{(k)}$ and bus voltages are contingency recourse. Thus (30) tests one shared active plan rather than separately optimized post-contingency active schedules.

The evaluation follows this chain. Experiments 1–2 test the incentive model and workload-state verifier; Experiments 3–5 isolate settlement design and network transfer; Experiments 6–7 test active workload constraints and participant allocation. Experiment 8 evaluates all RTS-24 line outages, Experiment 9 tests the telemetry-robust payment certificate and freezes its candidate target and power scale on validation data, Experiment 10 transfers the counterfactuals to nonlinear AC models, and Experiment 11 exhausts all declared spatial assignments and power scales. Experiment 12 embeds each submitted event

trajectory in a continuous real-arrival horizon and evaluates complete-cycle remuneration and bilateral participation. Experiment 13 replays the measured MIT/DCGM batch trace without intervening in the event, Experiment 14 solves one exact job-level flow LP for all joined jobs, and Experiment 15 independently audits the locked payment profile at the held-out q_{01} and q_{99} conversion endpoints. Experiment 16 audits immutable ledger provenance and reconciles the regional capacity with the same-source execution envelope. Experiment 17 fixes the information set at the event gate and compares a decision-time ledger with the complete-ledger audit. Experiment 18 evaluates AC post-contingency recourse across RTS-24, IEEE-30, IEEE-39, and IEEE-118 for every outage in a deterministic native-case admissibility set. Each confirmatory comparison uses the same 54 locked days unless a fixed network audit rule is stated; the case study is explanatory.

V. RESULTS

A. Incentive boundary and independent verification

Experiment 1 solves the entire 8×7 DR-price/call-probability grid. The dual threshold (9) agrees with the independently solved strategic program in all 56 cells. At \$150/MWh and $q = 0.35$, strategic reference scheduling inflates the event baseline by 4.09 MWh. Of 52.46 MWh credited, 18.05 MWh is unsupported by physical reduction, giving a false-credit ratio of 0.344. The profitable region follows the derived marginal boundary rather than a fitted classifier; the complete grid and its phase diagram are retained in the experiment output.

Experiment 2 evaluates all twelve methods on 54 later days. Table II separates the high-recall tail-risk counterfactual from the settlement-safe output. The tail-risk counterfactual improves mean credit F_1 from 0.485 to 0.513 over the selected single projection (Holm-adjusted $p = 0.0001$); its nRMSE is 0.353 versus 0.350, and the block test does not reject equality. The risk-constrained verifier combines seven independently feasible candidates without degenerating to one candidate. Relative to the equally informed feasible-quantile projection, unsupported credit falls from 2.026218 to 1.910590 MWh/day, a 5.71% reduction (paired three-day moving-block 95% interval $[0.057200, 0.534363]$ MWh; Holm-adjusted $p = 0.015625$), while nRMSE falls from 0.353737 to 0.348425 (1.51%, interval $[0.001547, 0.011528]$, Holm-adjusted $p = 0.010468$). The F_1 change is 0.485947 to 0.482752 (interval $[-0.004470, 0.002890]$, Holm-adjusted $p = 1.000000$), so the result is a selective reduction in unsupported credit rather than a claim of uniformly better classification. The raw quantile learner exposes 87.223 MWh, showing that the principal gain is the auditable workload-feasible set and its deterministic credit envelope, not privileged execution truth. The same-ledger literature panel reports exact implementations of incentive-compatible spatial DR, price-responsive non-wire scheduling, proactive workload shifting, and aggregator coordination; its four 54-day means are retained in the baseline ledger rather than conflated with the proposed selector.

The new risk fit is a convex quadratic program rather than a nonsmooth sorted-positive-part optimization: the selected reserve is $\beta = 1$ after four contiguous validation folds, and the final fit has KKT stationarity 9.4×10^{-10} and zero primal residual. On the locked test days, the final total false-credit exposure is 412.687 MWh-slot versus 422.677 for the single-projection reference, and its 75% daily-tail exposure is 26.051 versus 26.617 MWh-slot; both aggregate certificates are satisfied. The two-sided certificate in (19) is checked independently for all 54 locked days and every event cell. With the predeclared 1-MW band, the mean false-credit exposure is

TABLE II
LOCKED 54-DAY COUNTERFACTUAL RESULTS (DAILY MEANS).

Method	nRMSE	False MWh	Under-credit	F_1
High-5-of-10	2.653	118.820	7.496	0.230
Ridge	1.845	90.582	4.820	0.321
Gradient boosting	2.192	103.378	7.434	0.256
Extra Trees	2.173	100.920	7.247	0.263
Online metadata	2.123	98.681	6.795	0.275
Post-event metadata	1.968	91.107	6.610	0.299
Post-event quantile	1.891	87.223	6.579	0.324
Synthetic control	1.893	88.185	9.163	0.270
Feasible quantile	0.354	2.026	15.225	0.486
Single feasible	0.350	1.957	15.245	0.485
Tail-risk feasible	0.353	4.180	13.153	0.513
Risk-constrained safe	0.348	1.911	15.314	0.483




Fig. 3. Validation-fitted projection weights, constraint-set ablation, specification robustness, and the precision-recall frontier. Every point is produced by a globally solved workload program.

1.910590 MWh/day and the mean under-credit is 15.313968 MWh/day, compared with 15.244514 MWh/day for the selected single reference; every upper and lower pointwise margin is nonnegative. The RiskSafe therefore occupies a conservative point on an explicit false-credit/under-credit frontier: relative to feasible quantile it reduces false credit by 5.71% but increases under-credit by 0.089 MWh/day and does not improve F_1 . The pointwise envelope prevents arbitrary downward collapse relative to the selected single profile, but it does not make the accuracy tradeoff disappear. We report this cost directly instead of treating the certificate as a uniform predictive improvement.

The closest domain baselines are implemented on the same locked ledger rather than represented by generic machine-learning surrogates. Their daily means are retained in the complete baseline panel: incentive-compatible spatial DR (nRMSE 2.108, false credit 89.846 MWh, $F_1 = 0.206$), price-responsive non-wire scheduling (0.361, 2.319 MWh, 0.444), proactive workload shifting (0.350, 1.957 MWh, 0.485), and aggregator coordination (0.552, 0 MWh, 0). The incentive-compatible and flexibility formulations follow [5], [23]; the receding-horizon and aggregator implementations follow [6], [10]. The complete day-level rows are retained with the released experiment outputs.

Fig. 3 reports the complete day-level distributions, constraint ablation, validation frontier, and independent intervention robustness. The 5-by-3 capacity/deadline factorial contains 810 optimal schedules; maximum capacity- and deadline-binding shares are 0.310 and 0.948. The measured nonzero effects establish that the associated constraints are not decorative. Matrix nonzeros scale with log-log slope 1.003 over 96–608 intervals, matching the linear-horizon proposition.

TABLE III
BASE-CASE AND COMPLETE N-1 SETTLEMENT ERRORS (USD/DAY).

Panel	Settlement	Trace MAE	Workload MAE	Trace median/overpay	Workload median/overpay
Base case	Uniform gross	741.72	632.58	—	—
Base case	Nodal gross	741.70	632.11	—	—
Base case	Uniform signed	6.01	524.05	—	—
Base case	Nodal signed linear	6.00	524.23	—	—
Base case	Nodal exact net value	2.34	523.25	—	—
N-1	Base exact	192.69	207.67	98.50/185.21	126.36/156.89
N-1	N-1 linear	2.30	155.25	0.61/2.08	111.83/37.53
N-1	N-1 exact	0.90	155.40	0.82/0.01	112.62/36.97

B. Nodal value and spatial response

Experiment 3 crosses thirteen counterfactuals with five settlement mechanisms. Table III first uses the trace-anchored reference to isolate the mechanism. Exact signed nodal value has mean absolute error \$2.34, versus \$741.70 for nodal gross response. Under the workload verifier, the errors are \$523.25 and \$632.11, respectively; the remaining error is dominated by counterfactual uncertainty rather than value-function linearization. The paired one-factor decomposition prevents attributing this bundled improvement to location alone. With the trace reference, signed netting lowers error by \$735.69/day (Holm-adjusted exact block-sign $p = 3.1 \times 10^{-5}$), whereas changing uniform signed prices to nodal prices changes it by only \$0.0062/day ($p = 0.0874$), and replacing nodal linear value by exact value lowers it by \$3.66/day ($p = 0.0874$). Under the workload verifier, the corresponding effects are \$109.34, $-\$0.182$, and $\$0.988$ /day, none significant after Holm correction ($p = 0.706$). Thus signed rebound accounting identifies the dominant base-case gain; locational and security effects require the congested and N-1 panels below. All 5,616 interval/counterfactual instances satisfy (20); the largest numerical violation is 2.03×10^{-11} dollars.

Experiment 4 applies a predeclared rule—the largest gross-minus-net spatial offset—and selects day 69. The source curtailment, remote rebound, nodal prices, and line utilization are retained in the experiment output. The case explains why positive-only local accounting fails: migration creates a simultaneous negative component that gross payment discards.

C. Cross-network and N-1 security evidence

Experiment 5 contains 6,480 outcomes over four public networks, three native-load multipliers, two counterfactual qualities, five mechanisms, and 54 days. Congestion arises endogenously in three of 12 network-load cells. Exact value improves trace-anchored absolute error in all 12 cells (mean \$1.92/day; ten Holm-adjusted block tests below 0.05), but improves only five end-to-end workload cells (across-cell mean \$0.016/day). This separation supports Proposition 3 without converting a trace diagnostic into an end-to-end dominance claim.

The near-affine RTS-24, 0.98-load cell is recomputed on all 54 days at 10/40, 20/80, 40/160, and 80/320 settlement/evaluation segments. The four signed-linear-minus-exact differences are -0.01695 , -0.02730 , -0.01938 , and -0.01740 USD/day, all below 0.03 in magnitude, confirming local affinity rather than a discretization artifact.

Experiment 6 crosses capacity multipliers 0.55, 0.70, 0.85, 1.00, 1.20 and deadline multipliers 0.25, 0.50, 1.00 over 54 days (810 globally solved schedules). Capacity- and deadline-binding shares reach 0.0411 and 0.9118; mean false-credit exposure across capacity levels is 2.641, 2.225, 1.971, 1.832, and 1.941 MWh. Every schedule satisfies the declared workload constraints.

Experiment 7 enumerates all 16 coalitions for four participants and all 256 for eight. Trace-reference Shapley error is zero with maximum budget residual 2.3×10^{-13} dollars; workload-verifier error is 107.28 while its budget residual

is 1.6×10^{-13} . The exchangeable scaling panel reaches 20 participants (1,296 states; residual 2.98×10^{-13}), confirming exact budget balance rather than permutation sampling.

Experiment 8 enforces (25) for all 37 non-islanding outages of RTS-24 (the single islanding line is reported separately). Over 54 locked days and 2,592 interval-mechanism outcomes, native ratings, 0.90 load, 6% penetration, and four fixed sites are held constant; a 10-segment certificate is scored by an independent 40-segment N-1 dispatch. Trace-anchored mean errors for base-case exact, N-1 linear, and N-1 exact are \$192.69, \$2.30, and \$0.90 per day; the last improves on base-case exact by \$191.79/day (Holm-adjusted $p = 3.1 \times 10^{-5}$), whereas its \$1.40/day advantage over N-1 linear is not significant ($p = 0.122$). Workload-verifier errors are \$207.67, \$155.25, and \$155.40 (linear-to-exact $p = 0.690$), showing that counterfactual error dominates value linearization. Maximum post-contingency utilization is 1.000 p.u.; signed overpayment is retained in Table III.

Experiment 9 tests (27). All 54 daily lexicographic program pairs attain global optimality; the seven-profile candidate hull is anchored to the validation-selected single projection, with risk-constrained verifier external. The held-out factors 0.734, 0.996, and 1.278 form Ξ_3 ; all 162 certificates have zero positive cap violation and mean margins \$0.596, \$0.809, and \$1.038/day. A 384-cell validation-only N-1 panel selects $\rho = 1$ without locked-test payment or peak outcomes. Independent 40-segment errors are \$156.370/day for payment-certified verifier, \$155.605 for the selected single projection, \$155.703 for feasible quantile, and \$155.562 for risk-constrained verifier. These scores do not replace the deterministic cap. The maximum post-contingency loading remains 1.000 p.u. Experiment 15 separately audits the closed held-out conversion interval and is interpreted through Proposition 5. Its two-endpoint feasible payment segment has mean width \$0.545543/day at q_{01} and \$4.065188/day at q_{99} (maximum widths \$14.483473 and \$128.381395/day). This interval is an explicit uncertainty cost, not a lower-error claim. The cap is therefore a contractual upper bound under the declared market model; it is not evidence that the payment-certified profile has lower independent prediction error than every comparator.

Experiment 10 replaces the linear network with nonlinear AC OPF [12]. Across RTS-24, IEEE-30, IEEE-39, and IEEE-118, all 864 base outcomes converge with zero voltage violation and at most 1.000-p.u. loading; the payment guarantee remains scoped to (26). The corrective panel adds 5,400 OPFs over IEEE-9/14 non-islanding outages. A shared active plan is then fixed across all six IEEE-9 outages as in (30); its 3,888 cells converge with zero plan deviation and maximum loadings 0.876/0.899/0.923 p.u.

Experiment 11 removes fixed trace placement and power scale by crossing all $4! = 24$ region-to-bus mappings, 3%, 6%, and 9% peak penetrations, four counterfactuals, and all 54 days, yielding 15,552 outcomes. The risk-constrained verifier reduces mean absolute payment error relative to the closest feasible-quantile projection from \$95.475814/\$191.816652/\$289.187860 to \$91.261724/\$183.245675/\$276.253993 per day, and is lower in all 72 mapping-scale cells. The public IEEE-118 native load is fixed at the predeclared 0.90 multiplier used by the N-1 stress panel, while native branch ratings remain unchanged. Maximum loading and locational marginal price spread are 1.000 p.u. and \$4.06/MWh. The corresponding panels separate the conversion-scenario cap, shared-active-plan AC envelope, and spatial-scale transfer from the base-case accuracy results, so these claims are not conflated with the trace-anchored mechanism isolation

TABLE IV
CONTINUOUS ROLLING-HORIZON MARKET VALIDATION OVER 54 LOCKED DAYS.

Metric	payment-certified verifier	Quantile	risk-constrained verifier	Single
Complete-cycle MAE (\$/event)	0.084761	0.053170	0.241926	0.084761
Event-only MAE (\$/event)	8.091153	11.655725	11.186565	10.763272
Max window residual (MW)		6.0×10^{-6}		
Max cycle-energy residual (MWh)		3.0×10^{-6}		
Verified capacity (MWh)	13.902	13.902	13.902	13.902
Positive contracts		54/54		

TABLE V
INDEPENDENT REPLAY, INTERVAL-PAYMENT, AND LEDGER-PROVENANCE CERTIFICATES.

Evidence	Quantity	Baseline / value	Verified result
Trace replay	Mean slot MAE (MW)	Batch-only 6.615	Trace-constrained 4.287
Trace replay	Mean energy discrepancy (MWh)	Batch-only 635.058	Trace-constrained 411.516
Trace replay	Slot coverage	–	100% over 54 days
Flow witness	Jobs / service variables	–	71,128 / 1,974,690
Flow witness	Energy / maximum residual	–	2.657346 MWh / 9.4×10^{-18} MWh
Capacity	Observed peak / nameplate	–	117.936 / 118 MW = 0.9995
Interval certificate	Endpoint margins q_{01}/q_{99}	–	\$0.546 / \$4.065 per day
Interval certificate	Maximum cap violation	–	1.7×10^{-8} USD

result. Experiment 16 adds a direct calibration sensitivity: held-out measured-to-predicted energy ratios are 0.5845, 0.7342, 0.9957, 1.2780, and 2.6931 at the 1st, 10th, 50th, 90th, and 99th percentiles. Applied to the flexible MIT batch envelope, the direct q_{90} and q_{99} scales exceed the 118-MW commitment; the released sensitivity file reports the capacity-safe clipped factor (about 1.0005 at those endpoints) separately from the validation-frozen 3%–9% network panel. This is a calibration bound, not a claim that public traces identify a field geography. The complete assignment–penetration panel is retained in the experiment output.

Experiment 12 removes the daily reset: each event is embedded in 1,216 slots (608 pre-event, 96 event, 512 recovery), with only the eight event intervals projected. Across 216 method-day outcomes, maximum event L1 and cycle-energy residuals are 6.0×10^{-6} MW and 3.0×10^{-6} MWh. Complete-cycle MAE is \$0.084761/event for payment-certified verifier, \$0.053170 for feasible quantile, \$0.241926 for risk-constrained verifier, and \$0.084761 for single projection; event-only MAE is \$8.091153, \$11.655725, \$11.186565, and \$10.763272. Mean verified capacity is 13.902285 MWh at \$150/MWh; all 54 Nash contracts activate (mean surplus \$2,034.810894; minimum utility \$1.842918), with site residual below 1.42×10^{-14} . Table IV reports the cycle certificate.

D. Independent trace and deployment-boundary audits

The independent audit separates measured execution from the semi-synthetic intervention. Replaying immutable MIT scheduler/DCGM intervals over 54 days gives slot MAE 4.287 versus 6.615 MW for the batch-only trace and 100% slot coverage; this is observational fidelity, not a causal event. A sparse flow LP retains all 71,128 positive-energy jobs (1,974,690 service variables), reconstructs 2.657346 MWh, and has residuals below 10^{-17} MWh. The observed regional peak is 117.936 MW against the precommitted 118-MW nameplate. These results establish ledger fidelity, not semantic-field authenticity.

The held-out conversion-interval audit evaluates $q_{01} = 0.5845$ and $q_{99} = 2.6931$ with an independent 40-segment evaluator. Endpoint margins are \$0.545543/day and \$4.065188/day, with maximum cap violation 1.7×10^{-8} USD; Proposition 5 therefore certifies only the worst-case cost over the declared closed interval. The two-endpoint feasible payment segment has mean widths of \$0.545543/day and \$4.065188/day at these endpoints, respectively. The decision-time audit removes all post-gate arrivals: the truncated-ledger diagnostic changes nRMSE from 0.348425 to 0.525448 and credit recall from

0.389559 to 0.041674, whereas the committed-ledger pointwise-safe mode has zero false-credit exposure and a zero event-credit recall by design because it does not credit uncommitted future work. Finally, 1,698 declared AC corrective recourse rows are solved across four public networks and three validation-scaled penetrations. These are model and information-boundary audits, not a preventive AC security-constrained OPF certificate or a field-event result.

VI. CONCLUSION

This paper develops a provenance-bound workload-state verifier and a security-aware nodal settlement for spatially coupled AI data centers. The cumulative state enforces release, deadline, conservation, migration, and capacity constraints; the convex risk program uses an explicit linear false-credit/CVaR epigraph; the two-sided credit band limits unsupported credited response and bounds additional under-credit; and signed $N-1$ value accounts for rebound and congestion. Across 54 locked days, false-credit exposure falls from 2.026218 to 1.910590 MWh/day relative to the matched feasible-quantile reference, while nRMSE falls from 0.353737 to 0.348425. The same comparison increases under-credit from 15.225 to 15.314 MWh/day and does not improve F_1 , so the empirical result is a risk tradeoff rather than uniform predictive dominance. The trace-anchored mechanism-isolation panel reduces $N-1$ exact-settlement error from \$192.69 to \$0.90/day, whereas the end-to-end workload-verifier errors remain \$207.67, \$155.25, and \$155.40/day for the three settlement variants. The job-level witness preserves 71,128 measured jobs with machine-precision residuals, and the provenance certificate binds the witness to the immutable scheduler/DCGM join and reconciles the precommitted 118-MW nameplate with the observed regional envelope. The committed-ledger event-gate mode has zero false-credit exposure but credits no uncommitted future work; the interval payment audit reports the resulting conversion uncertainty instead of hiding it. Held-out calibration sensitivity and the cross-network AC audit evaluate the declared capacity and topology boundaries. The public traces provide reproducible workload and network evidence rather than a utility field intervention. Deployment therefore requires a signed operator ledger, co-located facility telemetry, topology access, and a capacity-safe workload-to-power conversion; these requirements define the next intertemporal AC and utility-event validation stage.

REFERENCES

- [1] X. Wang and W. Tang, "Modeling and analysis of baseline manipulation in demand response programs," *IEEE Transactions on Smart Grid*, vol. 13, no. 2, pp. 1178–1186, 2022.
- [2] California Independent System Operator, "California iso baseline accuracy work group final proposal," tech. rep., California ISO, 2017.
- [3] Z. Liu, A. Wierman, Y. Chen, B. Razon, and N. Chen, "Data center demand response: Avoiding the coincident peak via workload shifting and local generation," *Performance Evaluation*, vol. 70, no. 10, pp. 770–791, 2013.
- [4] M. A. Adnan, R. Sugihara, Y. Ma, and R. Gupta, "Energy efficient geographical load balancing via dynamic deferral of workload," in *2012 IEEE Fifth International Conference on Cloud Computing*, pp. 188–195, 2012.
- [5] Y. Cao, M. Cheng, S. Zhang, H. Mao, P. Wang, C. Li, Y. Feng, and Z. Ding, "Data-driven flexibility assessment for internet data center towards periodic batch workloads," *Applied Energy*, vol. 324, p. 119665, 2022.
- [6] Y. Zhang, A. Tsigikaridis, I. C. Paschalidis, and A. K. Coskun, "Data center and load aggregator coordination towards electricity demand response," *Sustainable Computing: Informatics and Systems*, vol. 42, p. 100957, 2024.
- [7] Y. Zhou, F. Wei, S. Li, Z. Wang, J. Liu, and D. Yu, "Data center load modeling through optimal energy consumption characteristics: A path to simultaneously enhance energy efficiency and demand response quality," *Applied Energy*, vol. 393, p. 126095, 2025.
- [8] W. Zhang, L. A. Roald, A. A. Chien, J. R. Birge, and V. M. Zavala, "Flexibility from networks of data centers: A market clearing formulation with virtual links," *Electric Power Systems Research*, vol. 189, p. 106723, 2020.
- [9] W. Zhang and V. M. Zavala, "Remunerating space–time, load-shifting flexibility from data centers in electricity markets," *Applied Energy*, vol. 326, p. 119930, 2022.
- [10] W. Zhang, L. A. Roald, and V. M. Zavala, "Exploring the impacts of power grid signals on data center operations using a receding-horizon scheduling model," *IEEE Transactions on Sustainable Computing*, vol. 8, no. 2, pp. 245–256, 2023.
- [11] B. Stott, J. Jardim, and O. Alsac, "DC power flow revisited," *IEEE Transactions on Power Systems*, vol. 24, no. 3, pp. 1290–1300, 2009.
- [12] R. D. Zimmerman, C. E. Murillo-Sanchez, and R. J. Thomas, "MATPOWER: Steady-state operations, planning, and analysis tools for power systems research and education," *IEEE Transactions on Power Systems*, vol. 26, no. 1, pp. 12–19, 2011.
- [13] D. A. Tejada-Arango, P. Sánchez-Martín, and A. Ramos, "Security constrained unit commitment using line outage distribution factors," *IEEE Transactions on Power Systems*, vol. 33, no. 1, pp. 329–337, 2018.
- [14] National Institute of Standards and Technology and Q. Dang, "Secure hash standard," Tech. Rep. FIPS PUB 180-4, National Institute of Standards and Technology, 2015.
- [15] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge University Press, 2004.
- [16] A. E. Hoerl and R. W. Kennard, "Ridge regression: Biased estimation for nonorthogonal problems," *Technometrics*, vol. 12, no. 1, pp. 55–67, 1970.
- [17] J. H. Friedman, "Greedy function approximation: A gradient boosting machine," *The Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, 2001.
- [18] P. Geurts, D. Ernst, and L. Wehenkel, "Extremely randomized trees," *Machine Learning*, vol. 63, pp. 3–42, 2006.
- [19] A. Abadie, A. Diamond, and J. Hainmueller, "Synthetic control methods for comparative case studies: Estimating the effect of california's tobacco control program," *Journal of the American Statistical Association*, vol. 105, no. 490, pp. 493–505, 2010.
- [20] R. T. Rockafellar and S. Uryasev, "Optimization of conditional value-at-risk," *The Journal of Risk*, vol. 2, no. 3, pp. 21–41, 2000.
- [21] Q. Huangfu and J. A. J. Hall, "Parallelizing the dual revised simplex method," *Mathematical Programming Computation*, vol. 10, no. 1, pp. 119–142, 2018.
- [22] B. Satchidanandan, M. Roozbehani, and M. A. Dahleh, "A two-stage mechanism for demand response markets," *IEEE Control Systems Letters*, vol. 7, pp. 49–54, 2023.
- [23] M. Chen, C. Gao, M. Shahidehpour, and Z. Li, "Incentive-compatible demand response for spatially coupled internet data centers in electricity markets," *IEEE Transactions on Smart Grid*, vol. 12, no. 4, pp. 3056–3069, 2021.
- [24] J. John F. Nash, "The bargaining problem," *Econometrica*, vol. 18, no. 2, pp. 155–162, 1950.
- [25] North American Electric Reliability Corporation, "TPL-001-5.1: Transmission system planning performance requirements," tech. rep., North American Electric Reliability Corporation, 2020.
- [26] S. Samsi, M. L. Weiss, D. Bestor, B. Li, M. Jones, A. Reuther, et al., "The mit supercloud dataset," in *2021 IEEE High Performance Extreme Computing Conference*, pp. 1–8, 2021.
- [27] L. S. Shapley, "A value for n -person games," in *Contributions to the Theory of Games, Volume II* (H. W. Kuhn and A. W. Tucker, eds.), pp. 307–317, Princeton University Press, 1953.
- [28] Y. Wang, Y. Chen, Z. Li, X. Kang, Y. Fang, Y. Zhou, Y. Zheng, Z. Tang, X. He, R. Guo, X. Wang, Q. Wang, A. C. Zhou, and X. Chu, "BurstGPT: A real-world workload dataset to optimize large language model serving systems," in *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pp. 5831–5841, 2025.
- [29] S. Babaeinejadsarookolae, A. Birchfield, R. D. Christie, C. Coffrin, C. DeMarco, R. Diao, et al., "The power grid library for benchmarking ac optimal power flow algorithms," *arXiv preprint arXiv:1908.02788*, 2019.
- [30] C. Grigg, P. Wong, P. Albrecht, R. Allan, M. Bhavaraju, R. Billinton, Q. Chen, C. Fong, S. Haddad, S. Kuruganty, et al., "The IEEE reliability test system–1996," *IEEE Transactions on Power Systems*, vol. 14, no. 3, pp. 1010–1020, 1999.
- [31] H. R. Künsch, "The jackknife and the bootstrap for general stationary observations," *The Annals of Statistics*, vol. 17, no. 3, pp. 1217–1241, 1989.
- [32] S. Holm, "A simple sequentially rejective multiple test procedure," *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979.